

Supplementary material: Treatment allocation under spatial spillover

Andrew Walther, Ross Joseph Simpson, Jr., Ashkan Habib and Feng-Chang Lin

October 2, 2026

This supplement preserves full methods, results and supporting diagnostics from the long-form source used to derive the article. It distinguishes the full synthetic grid benchmark, the selected-surface allocation-risk pilot and the completed observed-incidence NC application. Queen and rook results, planning years and spillover regimes remain separately labeled. The application is a simulation-based design evaluation, not observed intervention efficacy. Restricted county data are excluded. Aggregate source tables, manifests and the offline visual companion are maintained in the public repository; the ethics/data-use statement remains an explicit submission placeholder in the article.

Contents

S1	Detailed simulation methods	3
	S1.1 Spatial structure and outcome model	3
	S1.2 Heterogeneous baseline incidence	3
	S1.3 Treatment assignment designs	4
	S1.4 Estimation	6
	S1.5 Simulation design, resampling, and metrics	7
S2	Detailed grid results and allocation-risk pilot	8
	S2.1 Design performance under queen contiguity	8
	S2.2 Formal statistical comparison	8
	S2.3 Design rankings by incidence configuration	11
	S2.4 Spatial dependence, spillover magnitude and regime	11
	S2.5 Rook contiguity	13
	S2.6 Non-oracle sensitivity analysis	15
	S2.7 Stability across treatment effect sizes	15
	S2.8 Robustness across blocks and the recommended design	17
	S2.9 SRS benchmark and allocation-specific risk	18
S3	Observed-incidence application: full methods and results	19
	S3.1 Purpose and observed inputs	19
	S3.2 Geographic adaptations and treatment budgets	19
	S3.3 Outcome model and Monte Carlo evaluation	22
	S3.4 Completed results: accuracy, inference and reach	22
	S3.5 Allocation downside and sensitivity	23
S4	The eight-design comparison	26
S5	Simulation mechanics	26

S6	Supplementary statistical comparisons	29
S7	Rook contiguity figures	30
S8	Oracle and non-oracle estimates	33
S9	Technical notes	33
	S9.1 Identification diagnostics	33
	S9.2 Estimation error and its dependence on τ	34
S10	Observed-incidence application detail	36

S1 Detailed simulation methods

S1.1 Spatial structure and outcome model

We simulated a 10×10 regular grid of 100 clusters under two definitions of adjacency. Under queen contiguity, two clusters are neighbors if they share an edge or a corner; under rook contiguity, they must share an edge. Queen contiguity is the primary analysis in the detailed analysis. Rook contiguity is a sensitivity analysis and the continuity case with earlier work, whose grids used rook neighbors. For each definition, the binary adjacency matrix was row-standardized to form the spatial weights matrix W , so that WZ gives, for each cluster, the proportion of its neighbors that are treated, and WY gives the average outcome among its neighbors. Row-standardization expresses each cluster’s spillover exposure as an average over its own neighborhood rather than a raw sum, which would otherwise depend on how many neighbors the cluster has. Within a scenario, the same neighbor type defines W in the outcome model, in the Isolation Buffer design, and in the fitted model.

Cluster-level outcomes Y were generated from a spatial lag model with a spillover term,

$$Y = (I - \rho W)^{-1} [\tau Z + S(Z) + \beta X + \epsilon], \quad \epsilon \sim N(0, \sigma^2 I), \quad (\text{S1})$$

where $Z \in \{0, 1\}^{100}$ is the treatment assignment produced by the design under evaluation, τ is the direct treatment effect of interest, $S(Z)$ is the spillover term defined in Equation S2, ρ is the spatial autoregressive coefficient governing dependence in the outcome beyond that induced by treatment and spillover, X is the baseline incidence covariate with coefficient $\beta = 1$, and $\sigma = 1$. The spatial multiplier $(I - \rho W)^{-1}$ propagates each cluster’s shocks through the network of neighbors. We tested $\rho \in \{0, 0.01, 0.20, 0.50\}$. When $\rho = 0$ the multiplier is the identity and the model is a linear regression with independent errors, $Y = \tau Z + S(Z) + \beta X + \epsilon$, in which spatial structure enters only through the spillover term; larger values of ρ induce progressively stronger spatial clustering in the outcome.

We evaluated two spillover regimes,

$$S(Z) = \gamma WZ \quad (\text{both arms}), \quad S(Z) = \gamma WZ \circ (1 - Z) \quad (\text{control only}), \quad (\text{S2})$$

where $\circ$ denotes the elementwise product. Under the control-only regime, only untreated clusters respond to the treatment status of their neighbors, which represents classical contamination: an untreated cluster’s outcome rises with the share of its neighbors that are treated. Under the both-arms regime, every cluster responds to its neighbors’ treatment regardless of its own assignment, so neighboring treatment adds to a treated cluster’s own effect. We tested four spillover magnitudes, $\gamma \in \{0.5, 0.6, 0.7, 0.8\}$, and five magnitudes of the direct effect, $\tau \in \{0.8, 1.0, 1.5, 2.0, 3.0\}$, with $\tau = 1$ as the primary scenario throughout.

S1.2 Heterogeneous baseline incidence

the earlier small-grid study [Walther et al., 2026] assumed a spatially uniform baseline outcome risk. Cluster-level health outcomes typically vary across space, and the designs studied here differ in whether and how they use that variation. We generated the baseline incidence covariate X , with values in $(0, 1)$, under three mechanisms. Under the independent (iid) mechanism, each X_i is drawn from a Uniform(0, 1) distribution, with no spatial structure. Under the spatially autocorrelated mechanism, $X = \Phi\{(I - \rho_X W_Q)^{-1} u\}$, where $u \sim N(0, I)$, Φ is the standard normal distribution function applied elementwise, W_Q is the queen weights matrix, and $\rho_X \in \{0.20, 0.50\}$. Under the Poisson mechanism, a spatially autocorrelated log-rate surface, $\log \lambda = \log(35/100,000) + (I -$

$\rho_X W_Q)^{-1}u$ with $\rho_X \in \{0.20, 0.50\}$, generates cluster death counts $C_i \sim \text{Poisson}(\lambda_i P)$, with an equal population of $P = 100,000$ in every cluster, so that a cluster at the base rate expects 35 deaths. The observed rates C_i/P are then rank-normalized to $X_i = (r_i - 0.5)/N$, where r_i is the rank of cluster i 's rate among the $N = 100$ clusters and tied rates share their average rank. Ties remain at this population size: the Poisson surfaces used here have between 60 and 76 distinct values among the 100 clusters. The base rate of 35 per 100,000 was chosen to be representative of published SUD incidence: it lies between the SUD incidence rates that Mirzaei and colleagues estimated for US men (49.3 per 100,000) and women (21.7 per 100,000) aged 20–64, close to their average [Mirzaei et al., 2019]. It is not the observed rate of the county data described in the Application section, which come from a different source, age range, period, and case definition. The incidence filters always use the queen weights W_Q , whichever neighbor type the scenario uses, so the incidence surfaces are identical under rook and queen contiguity. The iid mechanism and the two strengths of each of the other two mechanisms give five incidence configurations.

For each configuration we generated $K = 10$ incidence surfaces, $X_1, \dots, X_{10}$, each from its own seed. The same ten surfaces are used by every scenario in that configuration, across neighbor types, outcome-model parameters, designs, and values of τ ; the design draws and the outcome noise are shared more selectively, as described under Simulation design, resampling, and metrics. Each design is applied to the surface that generates the outcome: assignments for surface k are drawn from X_k , and the outcome and the fitted model for those assignments use X_k as well. This assumes that the incidence observed when a trial is planned is the incidence that drives its outcome (the matched-surface assumption). The target of each scenario is performance averaged over incidence surfaces; for example, $\text{MSE} = E_X E_{Z|X} E_\epsilon[(\hat{\tau} - \tau)^2]$, with X drawn from the configuration's generator, Z from the design given X , and $\epsilon \sim N(0, I)$. Bias, coverage, and power are defined in the same way.

S1.3 Treatment assignment designs

We implemented eight treatment assignment designs and present six of them as the primary comparison set in the detailed analysis. The other two, Saturation Quadrants and Balanced Halves, each share the mechanism of a retained design, and in the eight-design comparison at $\tau = 1$ neither could be distinguished from its counterpart under either neighbor type, by the Nemenyi and Holm-adjusted Wilcoxon signed-rank tests described under Simulation design, resampling, and metrics below, applied here to all eight designs. Saturation Quadrants assigns the saturations 20%, 40%, 60%, and 80% to the four quadrants in random order, without reference to incidence. Its mean squared error was 0.0900, against 0.0907 for Incidence-Guided Saturation Quadrants, under queen contiguity (Nemenyi $p = 0.98$; Holm-adjusted Wilcoxon signed-rank $p = 0.41$), and 0.0708 against 0.0717 under rook contiguity (Nemenyi $p = 1.00$; Wilcoxon $p = 0.82$). We retain one representative of the saturation mechanism, Incidence-Guided Saturation Quadrants. Saturation Quadrants had the marginally lower mean under both neighbor types and the better average rank in the eight-design comparison (2.100 against 2.350 under queen contiguity and 2.281 against 2.306 under rook contiguity), so this choice does not rest on performance; that ordering the saturations by incidence makes no detectable difference is itself a result of the detailed analysis. Balanced Halves divides clusters into two halves by incidence rank and treats 25 at random in each. Its mean squared error was 0.1299, against 0.1308 for Balanced Quartiles, under queen contiguity (Nemenyi $p = 1.00$; Wilcoxon $p = 0.88$), and 0.0842 against 0.0837 under rook contiguity (Nemenyi $p = 1.00$; Wilcoxon $p = 0.82$). The two designs differ only in how finely they stratify on incidence, and we retain the finer, four-stratum version. Balanced Halves had the marginally lower mean and the better average rank under queen contiguity (3.312 against 3.356), and Balanced Quartiles had both under rook contiguity (2.919 against 3.038), so this choice does not rest on performance either. Both related

variants, and the full eight-design comparison, are reported in Section S4.

Table S1 describes the six retained designs, and Figure S1 shows one assignment from each on a single illustrative incidence surface. Checkerboard is the strict block-stratified allocation of the earlier small-grid study, in which no two rook neighbors share an arm. the earlier small-grid study’s Block Stratified Sampling design also admitted relaxed allocations on its 3×3 grid, where the arms cannot be balanced; on its 2×4 and 3×4 grids, as on the 10×10 grid here, the strict allocations are the two checkerboard patterns. We use the two names interchangeably for this allocation and use Checkerboard in the rest of the detailed analysis. Checkerboard is the only deterministic design: it gives the same assignment in every draw, although each of its fits still receives its own outcome noise. Every other design was drawn independently 25 times for each incidence surface. Designs that rank clusters by incidence break ties at random, afresh for each draw. High Incidence Focus is random only through that tie-breaking; under the iid and spatial mechanisms incidence values are distinct, so it gives the same assignment in all 25 draws on a surface, and its assignment changes only from one surface to the next. Isolation Buffer builds its assignment from the neighbor type of the scenario, and it treats 38 clusters on average under rook contiguity and 21 under queen contiguity. Every other design treats exactly 50.




Figure S1: One sample assignment from each of the six designs on the 10×10 grid, with panels in the order of Table S1. All six assignments use a single illustrative incidence surface of 100 independent $\text{Uniform}(0, 1)$ values generated with a fixed seed; it is not one of the simulation’s incidence surfaces. Tile fill shows baseline incidence (darker is higher), a ring marks a treated cluster, and an X marks a control. Isolation Buffer is shown under queen contiguity. The Checkerboard panel is labeled Block Stratified Sampling (Checkerboard), the the earlier small-grid study name for this allocation (see Treatment assignment designs).

Table S1: The six treatment assignment designs compared in the detailed analysis, ordered by mean squared error under queen contiguity at $\tau = 1$, pooled over the five incidence configurations (Table S2). A cluster’s incidence rank is its rank by the baseline incidence surface X_k , with ties broken at random in every draw. Clusters treated is the number of clusters treated per assignment; for Isolation Buffer, which depends on the neighbor type, it is the mean over its assignments under each neighbor type.

Design	Assignment rule	Clusters treated	Deterministic
Incidence-Guided Saturation Quadrants	The grid is divided into its four 5×5 coordinate quadrants. Quadrants are ranked by mean incidence and given saturations of 80%, 60%, 40%, and 20% from highest to lowest mean; within each quadrant, that share of its 25 clusters is treated at random	50	No
Balanced Quartiles	Clusters are divided by incidence rank into four strata of 25. Twelve clusters are treated at random in each stratum, and one more in each of two randomly chosen strata	50	No
Isolation Buffer	Clusters are treated one at a time, each chosen at random from those that are neither treated nor adjacent to a treated cluster, until none remain (a greedy random maximal independent set); no two treated clusters are neighbors	38 rook, 21 queen (mean)	No
High Incidence Focus	The 50 clusters with the highest incidence rank are treated	50	No
2×2 Blocking	The grid is divided into 25 non-overlapping 2×2 blocks; two of the four clusters in each block are treated at random	50	No
Checkerboard (strict Block Stratified Sampling allocation in the earlier small-grid study)	Clusters whose grid coordinates (x, y) have an odd sum $x + y$ are treated, so treatment alternates from cluster to cluster; under rook contiguity every neighbor of a treated cluster is a control, and every neighbor of a control is treated	50	Yes

S1.4 Estimation

We estimated τ with two maximum likelihood spatial lag models, both fit to every simulated outcome and both including an intercept. The primary, oracle model has the spatial lag ρWY and the covariates Z , $\text{Spill} = S(Z)$, and X_k , so it matches the outcome model in Equation S1, including the true regime-specific spillover covariate from Equation S2. The oracle model assumes that the analyst knows the form of the spillover mechanism and has the exact exposure term, so its results are a best case for each design; we return to this assumption as a limitation. The non-oracle model, a sensitivity analysis, has the covariates Z and X_k only and omits the spillover term. Both models are fit to the same outcomes.

We fit both models with our own implementation of the maximum likelihood estimator of the

spatial lag model [Bivand and Piras, 2015], the same estimator that the `lagsarlm` function of the R package `spatialreg` computes with `method = "eigen"`. It maximizes the concentrated log-likelihood over ρ , computing the log-determinant from the eigenvalues of W [Ord, 1975], estimates the remaining coefficients by least squares given $\hat{\rho}$, and computes standard errors from the analytic information matrix. On 5,120 validation fits it agreed with `lagsarlm` to within 9.7×10^{-8} in $\hat{\tau}$ and to a relative difference of 1.1×10^{-8} in its standard error, at roughly 1% of the computing time; we also refit a 1% sample of the full simulation with `lagsarlm` itself as a check. Confidence intervals are 95% Wald intervals, $\hat{\tau} \pm 1.96 \widehat{SE}$.

Checkerboard illustrates why interspersing alone cannot guarantee precision. Under rook contiguity on the rectangular grid, $WZ = 1 - Z$, so treatment and spillover cannot be separately identified; the fitted coefficient on treatment estimates $\tau - \gamma$. Under queen contiguity, treatment is identified but control-only spillover is weakly separated from treatment. We retain Checkerboard in the main design comparison and flag aliased estimates. Exact matrix identities, boundary information and the rare Isolation Buffer aliases are detailed in Section S9. These diagnostics concern the models implemented here.

S1.5 Simulation design, resampling, and metrics

The full simulation crossed five values of τ , five incidence configurations, two neighbor types, four values of ρ , four values of γ , two spillover regimes, and eight designs, for $5 \times 5 \times 2 \times 4 \times 4 \times 2 \times 8 = 12,800$ scenarios, of which the six designs reported in the detailed analysis account for 9,600. We call a combination of incidence configuration, neighbor type, ρ , γ , and spillover regime a block; there are $5 \times 2 \times 4 \times 4 \times 2 = 320$ blocks, 160 for each neighbor type. At the primary scenario, $\tau = 1$, the six designs give 1,920 scenarios, and every block contains all six.

Within each scenario, we drew $J = 25$ assignments from each of the $K = 10$ incidence surfaces, for 250 fits per scenario, and gave every fit its own noise vector, so that fit (k, j) has outcome $Y_{kj} = (I - \rho W)^{-1}[\tau Z_{kj} + S(Z_{kj}) + \beta X_k + \epsilon_{kj}]$. Checkerboard's 250 fits share one assignment but have 250 distinct noise vectors. Given a surface, its 25 pairs of assignment and noise are independent, and the ten surfaces are independent of one another, so the surface is the primary sampling unit. Each estimator was fit 3.2 million times. SRS was subsequently run with the same matched surfaces, noise keys and 250-fit structure, bringing the complete benchmark to 14,400 scenarios and 3.6 million fits per estimator; the original eight-design results were unchanged. SRS assigns exactly 50 of the 100 clusters to treatment by complete randomization, ignoring incidence and geography.

Every random draw was seeded from a hash of a key that names the draw: the configuration and surface index for an incidence surface; the block, surface, and design for a set of 25 assignments; and the block for the noise. Each draw therefore depends only on its key, and the results are identical whether scenarios are run sequentially, in parallel, or in any order. The keys also give common random numbers. Within a block, all designs share the same noise vectors and incidence surfaces; the five values of τ reuse the same assignments and noise; and the oracle and non-oracle models are fit to the same outcomes. Assignments and noise are independent across blocks, but blocks within an incidence configuration share its ten surfaces.

For each scenario and estimator, we computed five metrics over the valid fits, those with a finite estimate and standard error: bias, the mean of $\hat{\tau} - \tau$; the standard deviation (SD) of $\hat{\tau}$; mean squared error, $MSE = \text{mean}\{(\hat{\tau} - \tau)^2\}$; coverage, the fraction of fits with $L \leq \tau \leq U$, where $L, U = \hat{\tau} \mp 1.96 \widehat{SE}$; and power, the fraction of fits with $L > 0$, a one-sided test of $\tau = 0$ at level 0.025. For both estimators, every one of the 250 fits in every scenario was valid, and the fitting routine raised no warning other than the aliasing flags described above.

Monte Carlo standard errors come from the surface-level metrics. For a metric m , we computed m_k on the 25 fits from surface k and set $\text{SE}(m) = \text{sd}(m_1, \dots, m_{10})/\sqrt{10}$, with intervals $m \pm t_{0.975, 9} \text{SE}(m)$. A difference between two designs is paired at the surface level: its standard error is that of $m_{d,k} - m_{d',k}$ over k . For a metric pooled over the five incidence configurations, we apply the same formula to the 50 configuration-by-surface means, each averaged over that surface’s 32 blocks; these 50 means are independent of one another. Section S5 gives the parameter grid, the seed keys, and the exact formulas for these metrics and standard errors.

We compared designs within each neighbor type with a Friedman rank-sum test on scenario MSE [Friedman, 1937], treating the 160 blocks as repeated measures across the six designs. At $\tau = 1$ we followed it with Nemenyi post-hoc comparisons [Nemenyi, 1963, Demšar, 2006] and paired Wilcoxon signed-rank tests [Wilcoxon, 1945] with Holm’s adjustment for multiple comparisons [Holm, 1979] for all 15 pairs of designs, and we repeated the Friedman test at each other value of τ . Blocks within an incidence configuration share their surfaces, so these tests treat dependent blocks as independent, which overstates the information about variation from surface to surface. For each pair of designs adjacent in the mean-MSE order, we therefore also ran a paired Wilcoxon signed-rank test over 50 independent units, the 5 configurations times 10 surfaces, with each unit’s MSE averaged over that surface’s blocks, and a t_{49} interval for the mean difference. Because the values of τ share assignments and noise, Friedman tests at different values of τ are not independent confirmations of one another.

We report results separately for each neighbor type, with queen contiguity primary, and no number is averaged over the two. Unless stated otherwise, every number in the Results is pooled over the five incidence configurations, and every table and figure caption states whether it pools the configurations or shows them separately. We report the ranking within each configuration separately and show the split wherever a conclusion changes.

S2 Detailed grid results and allocation-risk pilot

S2.1 Design performance under queen contiguity

At the primary scenario, $\tau = 1$, under queen contiguity and pooled over the five incidence configurations, Incidence-Guided Saturation Quadrants had the lowest mean squared error ($\text{MSE} = 0.091$), followed by Balanced Quartiles (0.131), Isolation Buffer (0.160), High Incidence Focus (0.245), 2×2 Blocking (0.319), and Checkerboard (1.087) (Table S2, Figure S2). Coverage of the 95% intervals was between 0.936 and 0.941 for all six designs (Figure S3). Power to reject $\tau = 0$ was 0.930 for Incidence-Guided Saturation Quadrants, 0.827 for Balanced Quartiles, 0.731 for Isolation Buffer, 0.642 for 2×2 Blocking, 0.610 for High Incidence Focus, and 0.547 for Checkerboard.

Incidence-Guided Saturation Quadrants also led block by block. Of the 160 queen blocks at $\tau = 1$, Incidence-Guided Saturation Quadrants had the lowest MSE in 95 (59%), Balanced Quartiles in 58 (36%), and Isolation Buffer in 7 (4%); no other design had the lowest MSE in any block.

S2.2 Formal statistical comparison

The Friedman test rejected equal MSE ranks across the six designs under queen contiguity at $\tau = 1$ ($\chi^2 = 451.57$, $\text{df} = 5$, 160 blocks, $p < 2.2 \times 10^{-16}$). The average within-block ranks were 1.43 for Incidence-Guided Saturation Quadrants, 2.14 for Balanced Quartiles, 3.69 for Isolation Buffer, 4.24 for 2×2 Blocking, 4.74 for Checkerboard, and 4.76 for High Incidence Focus. Nemenyi comparisons (critical difference 0.596) separated 11 of the 15 pairs (Section S6 shows the critical-difference diagrams). Incidence-Guided Saturation Quadrants ranked significantly better than every other

Table S2: Performance of the six designs at the primary scenario, $\tau = 1$, under queen contiguity (primary) and rook contiguity (sensitivity), pooled over the five incidence configurations. Each entry averages the 160 scenarios of one design and neighbor type (5 incidence configurations $\times$ 4 values of $\rho \times 4$ values of $\gamma \times 2$ spillover regimes), each summarizing 250 fits of the oracle model. Rows in both panels are in order of mean squared error (MSE) under queen contiguity. Monte Carlo standard errors of MSE, in parentheses, are computed from the 50 configuration-by-surface means (Methods); those of bias, coverage, and power are all below 0.007. Average rank is the mean of each design’s within-block MSE rank among the six designs over the 160 blocks (1 = lowest MSE). Clusters treated is the mean number of clusters treated per assignment. Under rook contiguity, τ is not identified for Checkerboard ($WZ = 1 - Z$); its estimates are kept and flagged, and its coefficient on Z estimates $\tau - \gamma$.

Design	MSE (SE)	Bias	Coverage	Power	Average rank	Clusters treated
<i>Queen contiguity (primary)</i>						
Incidence-Guided	0.091 (0.0008)	0.032	0.937	0.930	1.425	50
Saturation Quadrants						
Balanced Quartiles	0.131 (0.0011)	0.017	0.936	0.827	2.138	50
Isolation Buffer	0.160 (0.0011)	0.003	0.939	0.731	3.694	21.4
High Incidence Focus	0.245 (0.0047)	0.031	0.938	0.610	4.763	50
2 $\times$ 2 Blocking	0.319 (0.0032)	0.010	0.939	0.642	4.244	50
Checkerboard	1.087 (0.0101)	0.001	0.941	0.547	4.737	50
<i>Rook contiguity (sensitivity)</i>						
Incidence-Guided	0.072 (0.0006)	0.018	0.942	0.963	1.488	50
Saturation Quadrants						
Balanced Quartiles	0.084 (0.0006)	0.008	0.940	0.917	1.775	50
Isolation Buffer	0.133 (0.0009)	0.003	0.939	0.806	3.381	38.2
High Incidence Focus	0.206 (0.0030)	0.018	0.939	0.645	4.969	50
2 $\times$ 2 Blocking	0.128 (0.0012)	0.004	0.938	0.818	3.400	50
Checkerboard (flagged)	0.480 (0.0015)	-0.652	0.150	0.414	5.987	50

design, including Balanced Quartiles ($p = 0.0086$), and Balanced Quartiles ranked significantly better than the remaining four. The four pairs that Nemenyi comparisons did not separate all lay among the bottom four designs: 2 $\times$ 2 Blocking against Isolation Buffer ($p = 0.090$), High Incidence Focus ($p = 0.130$), and Checkerboard ($p = 0.170$), and Checkerboard against High Incidence Focus ($p = 1.00$). Holm-adjusted Wilcoxon signed-rank tests separated all 15 pairs.

The bottom three designs were ordered differently by average rank and by mean MSE. By mean MSE the order was High Incidence Focus (0.245), 2 $\times$ 2 Blocking (0.319), and Checkerboard (1.087); by average rank it was 2 $\times$ 2 Blocking (4.24), Checkerboard (4.74), and High Incidence Focus (4.76). Both reversals have the same source: the MSEs of 2 $\times$ 2 Blocking and Checkerboard depended strongly on the spillover regime. Each had a lower MSE than High Incidence Focus in all 80 both-arms blocks (mean MSE 0.062 for each, against 0.162) and a higher MSE in all 80 control-only blocks (0.576 and 2.111, against 0.329; see the spillover-regime results below). Average rank counts each block once, so their wins under the both-arms regime offset their losses under the control-only regime; mean MSE weights those losses by their size.

Because blocks within an incidence configuration share their ten surfaces, we also compared each pair of designs adjacent in the mean-MSE order over the 50 independent configuration-by-surface units. Every adjacent pair differed significantly. The mean paired differences in MSE, with t_{49} 95%




Figure S2: Mean squared error (MSE) of the oracle estimate of τ by design at $\tau = 1$, under queen contiguity. Columns of panels correspond to the four levels of outcome spatial dependence ρ , rows to the both-arms (top) and control-only (bottom) spillover regimes, and bar colors to the spillover magnitude γ . Each bar overlays the MSEs of the five incidence configurations for one combination of design, ρ , γ , and regime: the bar's height is the largest of the five, and the horizontal lines within it mark the other four. Designs are ordered by pooled mean MSE (Table S2). The vertical scale is set separately for each row. Figure S12 in Section S7 shows rook contiguity.




Figure S3: Coverage of the 95% confidence intervals by design at $\tau = 1$, under queen contiguity. Each box summarizes the coverage of the 160 scenarios of one design (5 incidence configurations $\times$ 4 values of ρ $\times$ 4 values of γ $\times$ 2 spillover regimes), each estimated from 250 fits. Designs are ordered by median coverage, so the order differs from the MSE order used elsewhere. The dashed red line marks the nominal 0.95 level. Figure S13 in Section S7 shows rook contiguity.

intervals, were -0.0401 (-0.0431 , -0.0372) for Incidence-Guided Saturation Quadrants against Balanced Quartiles, -0.0293 (-0.0332 , -0.0255) for Balanced Quartiles against Isolation Buffer, -0.0853 (-0.0954 , -0.0752) for Isolation Buffer against High Incidence Focus, -0.0738 (-0.0858 , -0.0617) for High Incidence Focus against 2×2 Blocking, and -0.7677 (-0.7893 , -0.7460) for 2×2 Blocking against Checkerboard, with Wilcoxon $p \leq 4.9 \times 10^{-9}$ in every case. The better design of each pair had the lower MSE in all 50 units, except High Incidence Focus against 2×2 Blocking, where it did so in 48 (96%). This test follows the mean-MSE order. It separated High Incidence Focus from 2×2 Blocking, which the Nemenyi comparison of within-block ranks did not, because each unit’s MSE averages over both spillover regimes, and there the control-only losses of 2×2 Blocking outweigh its both-arms wins.

S2.3 Design rankings by incidence configuration

Table S3 gives each design’s MSE within each of the five incidence configurations. Under queen contiguity, the mean-MSE order within every configuration was the same as the pooled order: Incidence-Guided Saturation Quadrants, Balanced Quartiles, Isolation Buffer, High Incidence Focus, 2×2 Blocking, and Checkerboard. The configuration changed each design’s MSE only slightly: in this order, each design’s highest MSE across the five configurations was lower than the next design’s lowest. Incidence-Guided Saturation Quadrants had the lowest MSE in 18 to 21 of the 32 blocks in each configuration. Coverage did not depend on the configuration either: within each configuration, every design’s coverage, averaged over its 32 scenarios, was between 0.931 and 0.943. The mechanism that generated the incidence surface therefore did not change the ranking. This includes High Incidence Focus, the design that uses the incidence ranks most directly: its MSE was no higher under the two Poisson configurations than under the iid and spatial configurations.

S2.4 Spatial dependence, spillover magnitude and regime

Under queen contiguity, neither the strength of spatial dependence in the outcome nor the spillover magnitude changed the ranking. Across $\rho \in \{0, 0.01, 0.20, 0.50\}$, no design’s MSE changed by more than 0.02 (Table S4). Across $\gamma \in \{0.5, 0.6, 0.7, 0.8\}$, the MSE of every design except Checkerboard also varied by less than 0.02, and Checkerboard’s varied between 1.052 and 1.108 with no trend in γ (Table S5). Under rook contiguity, γ mattered for Checkerboard alone, for the reason given in the rook results below.

The spillover regime was the one factor that changed the queen ranking (Table S6, Figure S2). Under the both-arms regime, Balanced Quartiles had the lowest MSE (0.042), followed by Incidence-Guided Saturation Quadrants (0.046), 2×2 Blocking (0.062), Checkerboard (0.062), Isolation Buffer (0.161), and High Incidence Focus (0.162). Under the control-only regime, the order was Incidence-Guided Saturation Quadrants (0.135), Isolation Buffer (0.159), Balanced Quartiles (0.219), High Incidence Focus (0.329), 2×2 Blocking (0.576), and Checkerboard (2.111); at larger values of τ , Isolation Buffer matched or passed Incidence-Guided Saturation Quadrants under this regime (see Stability across treatment effect sizes). For Isolation Buffer, no treated cluster has a treated neighbor, so $WZ = 0$ at every treated cluster and $WZ \circ (1 - Z) = WZ$. The two regimes in Equation S2 therefore define the same outcome model, and so the same expected MSE; the observed values, 0.161 and 0.159, differ only by Monte Carlo error, because the two regimes use separate draws. Every other design had a higher MSE under the control-only regime, and the increase was largest for Checkerboard and 2×2 Blocking, which went from close to the best designs under the both-arms regime to the two worst under the control-only regime. Incidence-Guided Saturation Quadrants was first or second under both regimes, at this and every other value of τ , which matters

Table S3: Mean squared error by design within each of the five incidence configurations at $\tau = 1$, under queen contiguity (primary) and rook contiguity (sensitivity). Each entry averages the 32 scenarios of one design, neighbor type, and configuration (4 values of $\rho \times 4$ values of $\gamma \times 2$ spillover regimes), each summarizing 250 fits of the oracle model. The spatially autocorrelated and Poisson configurations are labeled by the strength of dependence in the incidence surface, ρ_X . Rows in both panels are ordered by pooled mean MSE under queen contiguity (Table S2). Under rook contiguity, τ is not identified for Checkerboard; its estimates are kept and flagged.

Design	iid	Spatial		Poisson	
		$\rho_X = 0.20$	$\rho_X = 0.50$	$\rho_X = 0.20$	$\rho_X = 0.50$
<i>Queen contiguity</i>					
Incidence-Guided Saturation Quadrants	0.093	0.087	0.092	0.090	0.092
Balanced Quartiles	0.130	0.130	0.131	0.128	0.135
Isolation Buffer	0.160	0.159	0.159	0.162	0.161
High Incidence Focus	0.245	0.267	0.245	0.232	0.238
2×2 Blocking	0.311	0.315	0.323	0.326	0.320
Checkerboard	1.091	1.068	1.086	1.112	1.077
<i>Rook contiguity</i>					
Incidence-Guided Saturation Quadrants	0.072	0.072	0.073	0.073	0.069
Balanced Quartiles	0.081	0.084	0.084	0.084	0.085
Isolation Buffer	0.131	0.131	0.132	0.136	0.134
High Incidence Focus	0.205	0.213	0.200	0.202	0.211
2×2 Blocking	0.129	0.125	0.129	0.130	0.129
Checkerboard (flagged)	0.480	0.482	0.474	0.481	0.482

Table S4: Mean squared error by design across the four levels of outcome spatial dependence ρ , under queen contiguity at $\tau = 1$, pooled over the five incidence configurations. Each entry averages 40 scenarios (5 incidence configurations $\times$ 4 values of $\gamma \times 2$ spillover regimes). Rows are ordered by pooled mean MSE under queen contiguity (Table S2).

Design	$\rho = 0$	$\rho = 0.01$	$\rho = 0.20$	$\rho = 0.50$
Incidence-Guided Saturation Quadrants	0.090	0.087	0.093	0.093
Balanced Quartiles	0.129	0.129	0.132	0.133
Isolation Buffer	0.161	0.157	0.158	0.165
High Incidence Focus	0.243	0.242	0.249	0.248
2×2 Blocking	0.318	0.320	0.314	0.324
Checkerboard	1.095	1.077	1.084	1.092

because a trial team may not know in advance which regime applies.

The increase under the control-only regime was variance, not bias. The oracle model includes the true regime-specific spillover covariate, so it has no omitted-variable bias. At $\tau = 1$, in every combination of design and regime under queen contiguity, the absolute bias was at most 0.048 and squared bias, averaged over scenarios, was at most 2.2% of MSE. Checkerboard showed the mechanism most clearly: under the control-only regime its spillover covariate is a linear combination of the intercept and Z at every interior cluster, so only the boundary clusters separate the spillover effect from the treatment effect (Methods). Checkerboard’s MSE rose from 0.062 under the both-arms regime to 2.111 under the control-only regime, while its bias stayed near zero (0.004) and its coverage stayed at 0.942: its intervals widened with its variance. Its pooled MSE of 1.087 was

Table S5: Mean squared error by design across the four levels of spillover magnitude γ at $\tau = 1$, under queen contiguity (primary) and rook contiguity (sensitivity), pooled over the five incidence configurations. Each entry averages 40 scenarios (5 incidence configurations $\times$ 4 values of $\rho \times 2$ spillover regimes). Rows in both panels are ordered by pooled mean MSE under queen contiguity (Table S2). Under rook contiguity, τ is not identified for Checkerboard; its estimates are kept and flagged.

Design	$\gamma = 0.5$	$\gamma = 0.6$	$\gamma = 0.7$	$\gamma = 0.8$
<i>Queen contiguity</i>				
Incidence-Guided Saturation Quadrants	0.092	0.089	0.092	0.089
Balanced Quartiles	0.132	0.128	0.132	0.131
Isolation Buffer	0.159	0.163	0.160	0.159
High Incidence Focus	0.244	0.249	0.248	0.241
2 $\times$ 2 Blocking	0.312	0.314	0.328	0.323
Checkerboard	1.108	1.052	1.086	1.101
<i>Rook contiguity</i>				
Incidence-Guided Saturation Quadrants	0.069	0.072	0.075	0.072
Balanced Quartiles	0.083	0.085	0.083	0.083
Isolation Buffer	0.131	0.133	0.134	0.133
High Incidence Focus	0.207	0.209	0.205	0.203
2 $\times$ 2 Blocking	0.129	0.129	0.127	0.128
Checkerboard (flagged)	0.295	0.403	0.538	0.684

Table S6: Mean squared error by design under each spillover regime at $\tau = 1$, under queen contiguity (primary) and rook contiguity (sensitivity), pooled over the five incidence configurations. Each entry averages 80 scenarios (5 incidence configurations $\times$ 4 values of $\rho \times 4$ values of γ). Rows are ordered by pooled mean MSE under queen contiguity (Table S2). Under rook contiguity, τ is not identified for Checkerboard; its estimates are kept and flagged.

Design	Queen		Rook	
	Both arms	Control only	Both arms	Control only
Incidence-Guided Saturation Quadrants	0.046	0.135	0.045	0.098
Balanced Quartiles	0.042	0.219	0.042	0.125
Isolation Buffer	0.161	0.159	0.131	0.134
High Incidence Focus	0.162	0.329	0.164	0.249
2 $\times$ 2 Blocking	0.062	0.576	0.051	0.206
Checkerboard (flagged under rook)	0.062	2.111	0.477	0.483

almost entirely variance: squared bias, averaged over scenarios, was 0.0045, less than 0.5% of MSE (Figure S4 shows the decomposition for the Poisson configuration with $\rho_X = 0.20$). 2 $\times$ 2 Blocking showed the same pattern on a smaller scale, with MSE rising from 0.062 to 0.576 and a bias of 0.016 under the control-only regime.

S2.5 Rook contiguity

Under rook contiguity at $\tau = 1$, pooled over the five incidence configurations, the mean-MSE order was Incidence-Guided Saturation Quadrants (0.072), Balanced Quartiles (0.084), 2 $\times$ 2 Blocking (0.128), Isolation Buffer (0.133), High Incidence Focus (0.206), and Checkerboard (0.480) (Table S2; Figure S12 in Section S7), and the same order held within each of the five incidence configurations




Figure S4: Decomposition of mean squared error into squared bias and variance by design at $\tau = 1$ for one incidence configuration, Poisson with $\rho_X = 0.20$, under queen contiguity. For each design, the two stacked components are the squared bias and the variance (SD^2) of $\hat{\tau}$, each averaged over the 32 scenarios of that configuration (4 values of $\rho \times 4$ values of $\gamma \times 2$ spillover regimes). Designs are ordered by pooled mean MSE (Table S2). Figure S14 in Section S7 shows rook contiguity.

(Table S3). Incidence-Guided Saturation Quadrants had the lowest MSE in 97 of the 160 rook blocks (61%) and Balanced Quartiles in 60 (38%). Coverage was between 0.938 and 0.942 for every design except Checkerboard, whose coverage was 0.150 (Figure S13 in Section S7); within single incidence configurations, the corresponding ranges were 0.935 to 0.946, and 0.14 to 0.15 for Checkerboard. The Friedman test rejected equal ranks ($\chi^2 = 703.76$, $df = 5$, 160 blocks, $p < 2.2 \times 10^{-16}$). Nemenyi comparisons separated 13 of 15 pairs; the two they did not separate were Incidence-Guided Saturation Quadrants against Balanced Quartiles ($p = 0.74$) and 2×2 Blocking against Isolation Buffer ($p = 1.00$). Holm-adjusted Wilcoxon tests separated every pair except 2×2 Blocking against Isolation Buffer. In the surface-level comparison, Incidence-Guided Saturation Quadrants had a lower MSE than Balanced Quartiles in 48 of the 50 units (mean difference -0.0121 ; -0.0139 , -0.0102 ; Wilcoxon $p = 9.9 \times 10^{-10}$). 2×2 Blocking and Isolation Buffer were close: 2×2 Blocking had the lower MSE in 36 of the 50 units (72%; mean difference -0.0048 ; -0.0078 , -0.0017 ; Wilcoxon $p = 0.0027$), but Isolation Buffer had the slightly better average rank (3.38 against 3.40), and the two designs traded places at some values of τ (see the results for τ below). Every other adjacent pair differed in all 50 units.

Checkerboard under rook contiguity is a different case: τ is not identified, and the coefficient on Z estimates $\tau - \gamma$ (Methods). The estimates behaved accordingly. At $\tau = 1$, Checkerboard's mean bias was -0.502 , -0.600 , -0.704 , and -0.802 at $\gamma = 0.5$, 0.6 , 0.7 , and 0.8 , while its mean standard deviation stayed between 0.20 and 0.21; squared bias, averaged over scenarios, accounted for 0.438 of its MSE of 0.480 (Figure S14 in Section S7 shows the decomposition for the Poisson configuration with $\rho_X = 0.20$). Its coverage fell from 0.316 at $\gamma = 0.5$ to 0.027 at $\gamma = 0.8$. Because $WZ = 1 - Z$, the two spillover covariates in Equation S2 coincide, $\gamma WZ \circ (1 - Z) = \gamma(1 - Z) = \gamma WZ$, so the two regimes define the same outcome model and the same expected MSE; the observed values, 0.477 under the both-arms regime and 0.483 under the control-only regime, differ only by Monte Carlo error, because the two regimes use separate noise draws. Under queen contiguity, by contrast, Checkerboard's coverage stayed between 0.939 and 0.943 at every level of ρ ; under rook contiguity it

was between 0.134 and 0.161 (Table S7). Rook Checkerboard estimates were precise but centered on the wrong quantity, the opposite of its imprecise but valid queen estimates. We keep these estimates in every ranking and flag them.

Table S7: Checkerboard’s 95% confidence interval coverage across the four levels of outcome spatial dependence ρ at $\tau = 1$, under queen contiguity (primary) and rook contiguity (sensitivity), pooled over the five incidence configurations. Each entry averages 40 scenarios (5 incidence configurations $\times$ 4 values of $\gamma \times 2$ spillover regimes). Under rook contiguity, τ is not identified for Checkerboard; its estimates are kept and flagged. The nominal level is 0.95.

Neighbor type	$\rho = 0$	$\rho = 0.01$	$\rho = 0.20$	$\rho = 0.50$
Queen	0.941	0.943	0.940	0.939
Rook	0.161	0.156	0.148	0.134

S2.6 Non-oracle sensitivity analysis


The non-oracle model, which omits the spillover term, was biased for every design (Table S16 in Section S8 gives both models’ MSE, bias, and coverage for each design). At $\tau = 1$, and setting aside Checkerboard under rook contiguity, its bias ranged from -0.12 to -0.40 and its coverage from 0.51 to 0.90 across designs and neighbor types. Under queen contiguity its bias ranged from -0.123 (Incidence-Guided Saturation Quadrants) to -0.244 (2×2 Blocking) and its coverage from 0.709 (2×2 Blocking) to 0.896 (High Incidence Focus); under rook contiguity the most biased design other than Checkerboard was Isolation Buffer (bias -0.400 , coverage 0.513). For Checkerboard under rook contiguity, the non-oracle estimates were identical to the oracle estimates (MSE 0.480, bias -0.652 , coverage 0.150), because once the aliased spillover column is dropped the two models are the same.

Omitting the spillover term also removes a column that can be nearly collinear with Z , and the non-oracle estimates were accordingly less variable for some designs. Under queen contiguity, the non-oracle MSE was lower than the oracle MSE for every design, most of all for Checkerboard (oracle 1.087, non-oracle 0.108) and 2×2 Blocking (oracle 0.319, non-oracle 0.121). Under rook contiguity it was higher for every design except 2×2 Blocking (0.128 under both models) and Checkerboard (identical). These results do not change the recommendation. Incidence-Guided Saturation Quadrants also had the lowest non-oracle MSE under both neighbor types (0.087 under queen contiguity and 0.086 under rook contiguity), and where a design’s non-oracle MSE was lower, the estimate was biased and its intervals under-covered, so the lower MSE does not support valid inference.

S2.7 Stability across treatment effect sizes

Under queen contiguity, each of the two design orders was the same at every tested value of τ (0.8, 1.0, 1.5, 2.0 and 3.0): the order by average within-block rank did not change across these values, and neither did the order by mean MSE, although the two orders differ from each other in the bottom three designs (see Formal statistical comparison; Figure S5a shows mean MSE). The Friedman statistic was 455.49, 451.57, 429.01, 411.19, and 376.79 at these five values ($p < 2.2 \times 10^{-16}$ at each). Under rook contiguity the order was also stable (Figure S5b) except for the close pair of 2×2 Blocking and Isolation Buffer: by average rank, 2×2 Blocking was ahead only at $\tau = 0.8$, and by mean MSE, Isolation Buffer was ahead only at $\tau = 3$. The five values of τ share their

assignments and noise, so these five comparisons are not independent confirmations of one another; they show that, under queen contiguity, the order did not change across the five values of τ we tested, given the same draws.



Rook: τ not identified for Checkerboard ($WZ = 1 - Z$); its estimates are kept but flagged.

Figure S5: Mean squared error by design across the five values of the treatment effect, $\tau \in \{0.8, 1.0, 1.5, 2.0, 3.0\}$, under (a) queen contiguity (primary) and (b) rook contiguity (sensitivity). Each point averages the 160 scenarios of one design at one value of τ , pooled over the five incidence configurations. Shaded bands show ± 1 times the average of the scenario-level Monte Carlo standard errors of MSE; they describe the precision of a single scenario’s MSE, not of the plotted average, whose standard error is smaller (Table S2). The shared legend lists the designs in order of pooled mean MSE under queen contiguity at $\tau = 1$. The vertical scales of the two panels differ. The five values of τ share their assignments and noise.

Within that stable order, the margins moved. Under queen contiguity, the average rank of Incidence-Guided Saturation Quadrants rose from 1.41 at $\tau = 0.8$ to 1.85 at $\tau = 3$, but Balanced Quartiles did not gain on it: Incidence-Guided Saturation Quadrants had the lower MSE of the two in 103, 102, 102, 102, and 101 of the 160 blocks at $\tau = 0.8, 1, 1.5, 2$, and 3 , and the average rank of Balanced Quartiles stayed between 2.13 and 2.14. Isolation Buffer gained instead, and its average rank fell from 3.71 to 3.31. The MSE of Isolation Buffer stayed at 0.160 at every value of τ , while that of Incidence-Guided Saturation Quadrants rose from 0.088 to 0.115 and that of Balanced Quartiles from 0.130 to 0.152, so the gap between Incidence-Guided Saturation Quadrants and Balanced Quartiles stayed between 0.037 and 0.041. Isolation Buffer’s gain came under the control-only regime, where Incidence-Guided Saturation Quadrants had a lower MSE than Isolation Buffer in 75, 73, 56, 40, and 12 of the 80 blocks at the five values of τ . The mean control-only MSE of Incidence-Guided Saturation Quadrants rose from 0.131 at $\tau = 0.8$ to 0.182 at $\tau = 3$, while that of Isolation Buffer stayed between 0.159 and 0.160 (Section S9). At $\tau = 2$ the two were level (0.1596 for Isolation Buffer against 0.1604), and at $\tau = 3$ Isolation Buffer had the lower control-only MSE (0.160 against 0.182). Incidence-Guided Saturation Quadrants remained second under the control-only regime at those two values, second to Balanced Quartiles under the both-arms regime at every value of τ , and lowest in pooled MSE at every value of τ . Its worst case across blocks

was the lowest of the six designs at $\tau \leq 1.5$, but at $\tau = 2$ and $\tau = 3$ that of Isolation Buffer was lower (0.207 against 0.210, and 0.208 against 0.238). The mean bias of Incidence-Guided Saturation Quadrants also grew with τ , from 0.029 at $\tau = 0.8$ to 0.047 at $\tau = 3$, but squared bias, averaged over scenarios, stayed at or below 3.0% of its MSE, so the rise in its MSE was mostly variance.

Under rook contiguity, Incidence-Guided Saturation Quadrants kept the lowest control-only MSE at every value of τ , although its margin over Isolation Buffer narrowed, from 0.095 against 0.134 at $\tau = 0.8$ to 0.131 against 0.136 at $\tau = 3$, where it had the lower MSE in 47 of the 80 control-only blocks. Its worst case across rook blocks was the lowest of the six designs at every value of τ except $\tau = 3$, where that of Isolation Buffer was lower (0.173 against 0.174).

Except for Checkerboard under rook contiguity, coverage stayed close to 0.94 at every value of τ . Pooled over the five incidence configurations, every design’s coverage was between 0.936 and 0.941 at every value of τ under queen contiguity, and every design except Checkerboard was between 0.937 and 0.942 at every value of τ under rook contiguity. The coverage of rook Checkerboard, for which τ is not identified, rose from 0.138 at $\tau = 0.8$ to 0.150, 0.198, 0.269, and 0.442 at $\tau = 1, 1.5, 2$, and 3 (Section S9). Section S6 shows coverage across τ for each design.

S2.8 Robustness across blocks and the recommended design

Table S8 gives the distribution of MSE across the 160 queen blocks at $\tau = 1$. Incidence-Guided Saturation Quadrants had the lowest median (0.080) and the lowest worst-case MSE of any design (0.169). Isolation Buffer had the narrowest range (0.128 to 0.206), because its MSE does not depend on the spillover regime, but at a higher level. Balanced Quartiles had a low median (0.115) and a worst case of 0.290. The distributions of 2×2 Blocking and Checkerboard reflected their dependence on the spillover regime: both had low first quartiles (0.062 and 0.063) and high upper quartiles (0.580 and 2.093). Checkerboard had the highest MSE in all 80 control-only blocks; in the 80 both-arms blocks, the highest MSE belonged to High Incidence Focus in 42 and to Isolation Buffer in 38. Under rook contiguity, Incidence-Guided Saturation Quadrants also had the lowest worst-case MSE (0.138).

Table S8: Distribution of block-level mean squared error by design across the 160 blocks under queen contiguity at $\tau = 1$, pooled over the five incidence configurations (32 blocks in each). At a fixed value of τ , each block contributes one scenario per design. Rows are sorted by median, which here coincides with the order of pooled mean MSE under queen contiguity (Table S2).

Design	Best	Q25	Median	Q75	Worst
Incidence-Guided Saturation Quadrants	0.037	0.045	0.080	0.135	0.169
Balanced Quartiles	0.034	0.042	0.115	0.221	0.290
Isolation Buffer	0.128	0.150	0.159	0.169	0.206
High Incidence Focus	0.139	0.160	0.234	0.322	0.424
2×2 Blocking	0.048	0.062	0.271	0.580	0.719
Checkerboard	0.051	0.063	0.899	2.093	2.843

Among the six designs in the detailed comparison, incidence-guided saturation combines low pooled MSE with comparatively favorable control-only performance. The SRS benchmark below changes the practical interpretation: recommend the saturation approach when control-only spillover is plausible, while retaining SRS as a competitive option under both-arms spillover. The near performance of plain saturation does not establish an additional benefit from ordering saturations by incidence. Scenario-level worst cases summarize variation across parameter settings; they are distinct from conditional risk across assignments within one setting.

S2.9 SRS benchmark and allocation-specific risk

At $\tau = 1$ under queen contiguity, pooled over the five incidence configurations and all spatial/spillover settings, guided saturation has mean MSE 0.0907 against SRS 0.1285; plain saturation has 0.0900. Table S9 separates the regimes and includes all eight strategies. Guided saturation’s control-only ratio to SRS is 0.631, whereas its both-arms ratio is 1.083. SRS, Balanced Quartiles and Balanced Halves have similar both-arms accuracy. Small mean differences should not be interpreted as equivalence, nor should pooling the regimes obscure the difference in design guidance. Paired comparisons over 50 independent configuration-by-surface units support the lower pooled MSE of both saturation designs relative to SRS; the six-design Friedman rankings elsewhere describe that subset only.

Table S9: Full grid benchmark at $\tau = 1$, queen contiguity. Each regime-specific entry averages 80 scenarios (five incidence configurations, four ρ and four γ values). MSE and coverage use the oracle fit. SRS is complete randomization of 50 clusters. No value pools queen and rook.

Design	Both arms		Control only	
	MSE	Coverage	MSE	Coverage
Checkerboard	0.0624	0.940	2.1114	0.942
High Incidence Focus	0.1617	0.941	0.3292	0.935
Saturation Quadrants	0.0442	0.942	0.1359	0.929
Isolation Buffer	0.1608	0.939	0.1595	0.939
2×2 Blocking	0.0623	0.937	0.5761	0.941
Balanced Quartiles	0.0421	0.940	0.2195	0.932
Balanced Halves	0.0420	0.939	0.2179	0.933
Incidence-Guided Saturation	0.0459	0.939	0.1355	0.934
SRS	0.0424	0.941	0.2146	0.936

To assess allocation risk, we additionally conducted a focused grid pilot with queen weights, $\tau = 1$, two fixed surfaces from each incidence configuration and the parameter corners $\rho \in \{0, 0.5\}$ and $\gamma \in \{0.5, 0.8\}$. For assignment a , its risk is $m(a) = E_{\epsilon}[(\hat{\tau} - \tau)^2 \mid a, X]$. Nested outcome replication estimates this conditional MSE before comparing its distribution across allocations. This is different from the distribution of squared errors in a single allocation or MSE across parameter scenarios. The pilot uses 100 assignment draws for stochastic designs and 100 outcomes per unique assignment, with a selected same-allocation refinement to 400 outcomes; singleton supports receive their own outcome replication.

Averaged over the 40 fixed-surface/corner units per regime in the pilot, guided saturation’s control-only mean MSE is 0.1409 against SRS 0.2172, estimated allocation q90 is 0.1747 against 0.2980, and estimated worst-decile mean is 0.1936 against 0.3486. Under both arms, the corresponding mean MSE is 0.0458 against 0.0428 and worst-decile mean 0.0582 against 0.0544. Plain saturation also performs favorably under control-only spillover. These selected-unit results support a qualified downside argument for saturation; they are not the full-study grid estimates. The focused refinement preserves the control-only direction but reveals substantial outcome-noise contributions to estimated allocation variation, especially under both arms. Finite-outcome quantiles and worst-decile membership remain uncertain. We therefore use the completed NC application below as the next setting-specific assessment rather than extrapolate exact tail protection from this pilot.

S3 Observed-incidence application: full methods and results

S3.1 Purpose and observed inputs

The application asks which allocation designs estimate a known education-intervention effect accurately on this geography, when some designs use historical SUD incidence to direct treatment. It supports planning an education intervention for healthcare professionals involved in SUD prevention. There are no observed intervention outcomes and no estimated relationship between local SUD incidence and education ability. Any eventual effect on SUD incidence requires a separate clinical evaluation. A real pilot might realize only one allocation per planning year, making both average estimation error and poor-allocation downside relevant.

The assignment units are 58 community college service-area clusters that partition North Carolina’s 100 counties under the study mapping (Figure S6). Northampton is assigned to Halifax Community College, Bertie to Martin Community College, and Hertford alone to Roanoke-Chowan Community College; these are analysis conventions, not exclusive service commitments. Observed county-year deaths and SEER populations aged 18–64 were aggregated to the same named clusters for each year, 2018–2021. The corrected numerator supplied by Habib includes heart-failure deaths formerly excluded as presumed non-sudden (A. Habib, personal communication, September 25, 2026). It totals 23,523 deaths over 25,594,321 person-years. The earlier numerator reproduces the totals behind Habib’s unpublished report [Habib, 2026]; the corrected numerator is the application input. We do not infer undocumented upstream exclusion rules or zip-code allocation from the received county panel.

Annual incidence is $100,000D_{iy}/P_{iy}$, and pooled incidence is $100,000\sum_y D_{iy}/\sum_y P_{iy}$, with the latter population sum interpreted as person-years. All 58 clusters are present in each year, and counts and populations sum back to statewide totals. We hold these observed annual surfaces fixed, generate no synthetic incidence surface and report years separately (Figure S7). Restricted county sources remain outside the public repository. Cluster-level summaries are authorized for reporting; the institutional ethics and data-use statement remains pending author-supplied documentation.

S3.2 Geographic adaptations and treatment budgets

Named queen weights cached on legal county boundaries define primary adjacency; rook weights are a sensitivity. Row order, binary adjacency and row-standardized weights are verified against the legal-boundary geometry before simulation. Shared legal boundaries across water can create adjacency; this is a prespecified geographic exposure proxy, not a measured professional communication network.

All eight designs and SRS are evaluated. Checkerboard becomes a deterministic greedy graph assignment: visit clusters by decreasing degree, choose the arm that conflicts least with assigned neighbors, break equal conflicts toward control, and complement the assignment if more than half are treated. This is an interspersed analog, not a literal rectangular checkerboard or an exact bipartite coloring. High Incidence Focus treats the 29 highest annual-rate clusters. Isolation Buffer draws a greedy random maximal independent set on the active graph. SRS treats a random 29 of 58 clusters. Balanced Quartiles and Balanced Halves stratify annual rate ranks and treat 29 overall; odd-size strata receive their necessary extra treatments at random. In Balanced Halves, one half receives 14 treatments and the other 15, with the extra-treatment half chosen randomly.

The two saturation designs use one frozen four-region partition for all years. The existing geographic k -means procedure generates 250 candidates, selecting the partition by compactness and balance in reference population, cluster count and county count, without consulting outcomes. Reference population is pooled person-years divided by four. Plain saturation randomly assigns




Figure S6: The assignment of North Carolina’s 100 counties to the 58 community college service areas used as clusters in the planned application. Each county is colored by its primary serving college, and each service area carries a single number, keyed to the legend. Hatching marks the two counties that also lie within Roanoke-Chowan Community College’s broader service network: Northampton County, assigned to Halifax Community College (24), and Bertie County, assigned to Martin Community College (30). Roanoke-Chowan Community College (41) serves Hertford County alone in this analysis. The map shows county boundaries and the county-to-college assignment only; it does not depend on the SUD data.

20/40/60/80% to these regions each draw. Guided saturation first ranks clusters by annual per-capita rates, averages their ranks equally within each region, ranks those four means, and assigns 80/60/40/20% from highest to lowest. Regional ties are broken randomly per allocation. Local treatment counts are rounded under the retained design rules. Per-capita cluster rates already account for population size; regional mean ranks answer an equally weighted cluster question, whereas pooled regional deaths/population weights larger clusters more heavily. Mean raw rates and pooled regional rates are evaluated as prespecified sensitivities rather than chosen by performance. The grid rule averages X directly: mean ranks align with its rank-scaled Poisson cases but are a deliberate application adaptation, rather than the exact rule for continuous grid incidence.

The grid’s 2×2 Blocking becomes spatial blocking on 14 irregular groups, formed once per year by k -means on standardized projected coordinates, incidence ranks and $\log(1+\text{population})$, then frozen. Half of each group, rounded, is treated randomly. This is not a set of literal square blocks. In each year one singleton group remains control under $\text{round}(0.5)=0$. Budgets differ across designs and are disclosed rather than forced into an unapproved equal-budget modification (Table S10).

Observed SUD incidence: four fixed yearly planning settings
58 community-college clusters; denominator is population aged 18–64




Figure S7: Observed annual SUD incidence per 100,000 adults aged 18–64 in the 58 service-area clusters. These four fixed planning inputs direct incidence-aware allocation; they are not simulated outcomes or measures of education ability.

Section S10 maps the frozen regions and example assignments.

Table S10: Primary queen application treatment budgets. Entries are mean treated-cluster counts / percent of annual population aged 18–64. Counts and population shares use the same sampled assignments; annual population diagnostics are recalculated for shared supports. Shares are geographic reach, not healthcare-professional enrollment.

Design	2018	2019	2020	2021
Graph Checkerboard	27.00 / 50.4	27.00 / 50.4	27.00 / 50.4	27.00 / 50.7
High Incidence Focus	29.00 / 22.2	29.00 / 23.5	29.00 / 23.3	29.00 / 22.4
Plain saturation	29.03 / 49.2	29.03 / 49.2	29.03 / 49.1	29.03 / 49.1
Isolation Buffer	15.89 / 26.7	15.89 / 26.7	15.89 / 26.7	15.89 / 26.7
Spatial blocking	28.00 / 50.0	31.00 / 55.1	27.00 / 48.0	29.00 / 50.3
Balanced Quartiles	29.00 / 48.5	29.00 / 49.7	29.00 / 50.6	29.00 / 49.8
Balanced Halves	29.00 / 50.0	29.00 / 50.5	29.00 / 49.2	29.00 / 49.4
Guided saturation	29.00 / 53.2	29.00 / 45.1	29.00 / 47.9	29.00 / 51.8
SRS	29.00 / 49.5	29.00 / 49.5	29.00 / 49.5	29.00 / 49.5

S3.3 Outcome model and Monte Carlo evaluation

For the primary education outcome, we set $\beta = 0$ in Equation S1, use the 58-cluster W , and fix $\tau = 1$, $\sigma = 1$:

$$Y = (I - \rho W)^{-1}[\tau Z + S(Z) + \epsilon], \quad \epsilon \sim N(0, I). \quad (\text{S3})$$

Observed incidence affects assignments but neither baseline education nor the constant structural intervention effect. We estimate τ using the validated ML engine with an intercept, Z and the true regime-specific $S(Z)$, without X . As in the grid model, τ is the coefficient before spatial propagation, not an average total impact after applying $(I - \rho W)^{-1}$.

Primary queen settings cross $\rho = 0, 0.2, 0.5$, $\gamma = 0.5, 0.8$, both regimes, four years and nine designs (432 reporting rows). A matched hypothetical baseline sensitivity adds $X = (\text{rank}_{\text{average}}(\text{rate}) - 0.5)/58$ with $\beta = 1$ to both the outcome baseline and the fitted covariates (432 rows). It changes baseline and adjustment together and is not an empirical incidence–education association. Primary-model rook corners use $\rho = 0, 0.5$ (288 rows); the two alternative regional summaries add 96 guided-saturation rows. The full application has 1,248 rows from 936 distinct computational sources.

We initially sample $J = 100$ assignments and simulate $R = 100$ independent outcomes per unique assignment. Proven singleton supports receive $R = 1,000$, not 100 copies masquerading as new allocations. An adaptive, prespecified replication schedule extends allocations or outcomes until relative MCSE(mean MSE) is at most 5% and MCSE(coverage) at most 0.01. Two singleton matched-sensitivity sources require $R = 2,000$; all stochastic main settings retain $J = 100, R = 100$. Keys seed individual allocation and outcome indices reproducibly. Content-based support keys allow equivalent rules to share computation, with annual treatment/population diagnostics recalculated. Primary Checkerboard, plain saturation, buffer and SRS performance reuse the same source across years when annual incidence does not enter their allocation or outcome model. These repeated values are not independent annual confirmation.

For each allocation a , nested outcomes estimate $m(a) = E_{\epsilon}[(\hat{\tau} - \tau)^2 \mid a, \text{observed inputs}]$. We report its empirical distribution, estimated q90 and estimated mean among the worst 10% sampled draws alongside unconditional mean MSE, bias and coverage. Duplicate allocations retain their sampled frequency but reuse one conditional estimate; joint MC uncertainty accounts for the resulting covariance. A separate extension at queen $\rho = 0, 0.5$, $\gamma = 0.8$, both regimes and all years/designs retains the same allocations and extends stochastic outcomes to $R = 400$, preserving singleton $R = 1,000$. Opposite outcome halves select and evaluate worst-decile allocations to diagnose selection noise. The extension is not a new allocation sample. Actual fitted-matrix rank, residual treatment information, aliases, warnings, invalid estimates/intervals and spatial-parameter boundary fits are recorded explicitly. Section S10 gives the reproducibility and precision details.

S3.4 Completed results: accuracy, inference and reach

Production generated 7,598,000 independent outcome fits. Every setting is complete and passes its prespecified mean-MSE/coverage precision gates: maximum relative MCSE(mean MSE) 0.04871 and maximum MCSE(coverage) 0.00892. There are no warnings, aliases, boundary fits or estimation failures. The tail extension adds 2,280,000 independent outcomes; copied main prefixes are counted once, giving 9,878,000 distinct outcomes across the two stages. Monte Carlo precision describes uncertainty in simulated performance; it does not certify interval calibration.

Figure S8 compares all nine designs by year and regime. Under control-only spillover, plain saturation has descriptive mean MSE 0.22329 against SRS 0.30761, a 27.4% reduction. Guided saturation reduces mean MSE by 21.1–29.9% across the four observed years (Table S11). Both

saturation designs have lower mean MSE than SRS in every primary parameter setting, including setting-specific approximate Monte Carlo intervals for their differences. Neither is uniquely best; incidence ordering has not established an accuracy benefit over plain saturation. High Incidence Focus is year-sensitive under control-only spillover and cannot be recommended uniformly.

Under both-arms spillover, SRS mean MSE is 0.07638, plain saturation 0.07518, and guided saturation ranges from 0.07428 to 0.07682 across years. SRS remains competitive with balancing, blocking and saturation. Small point differences do not establish equivalence or universal optimality. High Incidence Focus has low both-arms MSE in the primary model, but its allocation is deterministic for these fixed rates; this is conditional model-based accuracy, not evidence of randomization-based causal validity.

Mean annual coverage across primary settings is approximately 92.6–93.6%, below nominal 95%. The full year/design/regime bias and coverage table is in Section S10. No design’s passing precision gate repairs this finite-sample undercoverage. Graph Checkerboard’s treatment effect is identified on this irregular graph, but weak treatment–spillover separation makes its estimates imprecise; this differs from exact rook non-identification on the rectangular grid.

Population shares contextualize performance. High Incidence Focus treats 29 clusters but only about 22–24% of working-age population, while buffer treats a mean 15.89 clusters and about 27%. These population denominators measure geographic reach, not the number of healthcare professionals educated, intervention cost or participant-level precision. Spatial-block treatment counts are 28/31/27/29 by year; complete budgets and shares are reported in Table S10 and the appendix.

Table S11: Primary application yearly mean MSE, queen contiguity; equal weights over the six ρ/γ pairs. Plain saturation and SRS repeat shared primary sources across years.

Year	Both arms		SRS	Control only	
	SRS	Guided sat.		Plain sat.	Guided sat.
2018	0.07638	0.07428	0.30761	0.22329	0.24285
2019	0.07638	0.07628	0.30761	0.22329	0.21565
2020	0.07638	0.07682	0.30761	0.22329	0.22029
2021	0.07638	0.07450	0.30761	0.22329	0.23732

S3.5 Allocation downside and sensitivity

At the two refined queen ρ corners with $\gamma = 0.8$, under control-only spillover, plain saturation’s estimated q90 is 31.8–32.9% below SRS and its estimated worst-decile mean 34.6–37.8% below. Guided saturation’s corresponding reductions are 20.8–38.6% and 20.7–43.4%. The pattern persists when one outcome half selects the worst decile and the opposite half evaluates it. Figure S9 shows the $\rho = 0.5$ corner. This supports a descriptive downside benefit alongside lower average error, conditional on the studied model and sampled assignments.

Tail uncertainty remains material. With $J = 100$, the worst decile contains only ten sampled draws; extra outcomes reduce noise in $\hat{m}(a)$ but do not enlarge the allocation sample. All 112 non-singleton reporting rows flag uncertain tail-boundary membership. Under both arms, guided saturation’s half-sample worst-decile Jaccard overlap is 0–0.25 and approximately 61–86% of raw estimated-risk variance is attributable to outcome Monte Carlo noise. Plain saturation has similarly unstable identities. Small both-arms tail differences therefore do not establish superior downside protection. Singleton Checkerboard and High Incidence Focus have no allocation variation for these




Figure S8: Completed primary application: mean MSE by design, planning year and spillover regime, queen contiguity and education outcome. Each entry equally averages the six ρ/γ pairs. Both regime columns use a common logarithmic MSE scale. Shared annual sources are repeated reporting, not independent replications; pooled Monte Carlo intervals are not supplied.

inputs; their q90 and worst-decile mean equal conditional mean MSE, which does not imply low error or absence of outcome uncertainty.

In the matched hypothetical incidence-baseline/X-adjusted sensitivity, High Incidence Focus MSE rises to roughly 3.68–4.44 times its primary value under both arms and 1.42–2.52 times under control-only spillover. Treatment–incidence alignment leaves less residual treatment information after adjustment; this sensitivity changes baseline and adjustment jointly and does not imply that SUD incidence predicts education ability. Rook/queen comparisons matched at the same parameter corners preserve broad regime-specific conclusions. Alternative regional summaries change region ordering in some years; equal summaries with identical ordering/ties reuse the same distribution. We retain the prespecified mean-rank rule rather than select a summary retrospectively. Section




Figure S9: Allocation-specific conditional MSE in the refined primary queen setting $\rho = 0.5$, $\gamma = 0.8$, by year and spillover regime. Mean MSE, estimated allocation q90 and estimated worst-decile mean summarize the same $J = 100$ sampled assignments, with at least 400 independent outcomes per stochastic assignment or 1,000 for singleton supports. These finite-outcome tail estimates are descriptive, with uncertain membership; they are not quantiles of realized errors or a fresh allocation confirmation.

S10 reports the matched sensitivity table and regional ordering.

The application supports considering either saturation strategy when control-only spillover is plausible, and SRS alongside other randomized strategies when both-arm exposure is expected. A trial team should consider likely spillover, recruitment, cluster and participant budgets, achievable measurement precision and interval calibration before selecting one realized allocation. The findings do not establish intervention effectiveness or a final trial protocol.

S4 The eight-design comparison

Table S12 defines the two related variants excluded from the detailed six-design grid ranking but included in the main-text nine-design benchmark and application, in the format of Table S1, and Table S13 reports the comparison of all eight designs at the primary scenario, $\tau = 1$.

Table S12: The two related variants outside the detailed six-design grid ranking, in the format of Table S1. Saturation Quadrants shares its mechanism with Incidence-Guided Saturation Quadrants, and Balanced Halves with Balanced Quartiles (see Treatment assignment designs). A cluster’s incidence rank is its rank by the baseline incidence surface X_k , with ties broken at random in every draw. Clusters treated is the number of clusters treated per assignment.

Design	Assignment rule	Clusters treated	Deterministic
Saturation Quadrants	The grid is divided into its four 5×5 coordinate quadrants, which are given saturations of 20%, 40%, 60%, and 80% in an order drawn at random for each assignment, without reference to incidence; within each quadrant, that share of its 25 clusters is treated at random	50	No
Balanced Halves	Clusters are divided by incidence rank into two strata of 50; 25 clusters are treated at random in each stratum	50	No

The argument for leaving Saturation Quadrants and Balanced Halves out of the detailed grid analysis, with the Nemenyi and Wilcoxon results for each pair, is given under Treatment assignment designs in Methods; Table S13 supplies the rest of the eight-design comparison. Under both neighbor types, each excluded design is adjacent to its retained counterpart in mean MSE. The MSEs and coverages of the six retained designs are those of Table S2, given here to four and three decimal places.

S5 Simulation mechanics

This section gives the mechanical details behind Simulation design, resampling, and metrics (Methods): the parameter grid, the random-number streams, the Monte Carlo structure, and the exact formulas for the metrics and their Monte Carlo standard errors.

Parameter grid. Table S14 lists the parameters that were crossed to form the 12,800 scenarios and the settings that were held fixed.

Random-number streams. Every random draw in the simulation has its own seed, computed by hashing a key that names the draw. The key lists the draw’s type (incidence surface, assignments, or noise) and the parameters that identify it, with numeric values written to a fixed number of decimal places so that equal values always give the same key. The random-number generator and its methods for normal and discrete draws are also fixed explicitly, so the streams do not depend on software defaults. Table S15 lists the three kinds of draw, the parameters in their keys, and the scenarios that therefore share them. The consequences of this scheme, independence from the order of execution and common random numbers, are described in Methods.

Monte Carlo structure and checkpoints. Within a scenario, fit (k, j) , with $k = 1, \dots, 10$ and $j = 1, \dots, 25$, uses surface X_k , assignment Z_{kj} , and noise vector ϵ_{kj} , the (k, j) column of

Table S13: Performance of all eight designs at the primary scenario, $\tau = 1$, under queen contiguity (primary) and rook contiguity (sensitivity), pooled over the five incidence configurations. Each entry averages the 160 scenarios of one design and neighbor type (5 incidence configurations $\times$ 4 values of $\rho \times 4$ values of $\gamma \times 2$ spillover regimes), each summarizing 250 fits of the oracle model. Rows in both panels are in order of mean squared error (MSE) under queen contiguity. Retained marks the six designs compared in the detailed grid analysis of the detailed analysis. Average rank is the mean of each design’s within-block MSE rank among the eight designs over the 160 blocks (1 = lowest MSE), so it differs from the six-design average rank in Table S2. Under rook contiguity, τ is not identified for Checkerboard; its estimates are kept and flagged.

Design	Retained	MSE	Coverage	Average rank
<i>Queen contiguity (primary)</i>				
Saturation Quadrants	No	0.0900	0.936	2.100
Incidence-Guided Saturation Quadrants	Yes	0.0907	0.937	2.350
Balanced Halves	No	0.1299	0.936	3.312
Balanced Quartiles	Yes	0.1308	0.936	3.356
Isolation Buffer	Yes	0.1602	0.939	5.138
High Incidence Focus	Yes	0.2455	0.938	6.763
2 $\times$ 2 Blocking	Yes	0.3192	0.939	6.244
Checkerboard	Yes	1.0869	0.941	6.737
<i>Rook contiguity (sensitivity)</i>				
Saturation Quadrants	No	0.0708	0.940	2.281
Incidence-Guided Saturation Quadrants	Yes	0.0717	0.942	2.306
Balanced Halves	No	0.0842	0.940	3.038
Balanced Quartiles	Yes	0.0837	0.940	2.919
Isolation Buffer	Yes	0.1329	0.939	5.225
High Incidence Focus	Yes	0.2062	0.939	6.969
2 $\times$ 2 Blocking	Yes	0.1281	0.938	5.275
Checkerboard (flagged)	Yes	0.4799	0.150	7.987

the block’s 100×250 noise matrix (see Simulation design, resampling, and metrics). The unit of computation is a combination of incidence configuration, neighbor type, and ρ : 40 units, each covering 4 values of γ , 2 spillover regimes, 8 designs, 5 values of τ , 250 fits, and 2 estimators. The units ran in parallel, and the run stops if any unit fails. Each unit’s results are saved together with a record of the parameter settings, the simulation code, and the software versions that produced them, and saved results are reused only if all of these match the current run, so results from different code or settings are never combined.

Scenario metrics. For one scenario and estimator, let $f = 1, \dots, n$ index the valid fits, those with a finite estimate and standard error (here every fit was valid, so $n = 250$). Let $L_f, U_f = \hat{\tau}_f \mp z_{0.975} \widehat{SE}_f$, where $z_{0.975} = \Phi^{-1}(0.975) \approx 1.96$, and let $\bar{\tau} = n^{-1} \sum_f \hat{\tau}_f$. Then

$$\text{Bias} = \frac{1}{n} \sum_{f=1}^n (\hat{\tau}_f - \tau), \quad \text{SD} = \left\{ \frac{1}{n-1} \sum_{f=1}^n (\hat{\tau}_f - \bar{\tau})^2 \right\}^{1/2}, \quad \text{MSE} = \frac{1}{n} \sum_{f=1}^n (\hat{\tau}_f - \tau)^2,$$

$$\text{Coverage} = \frac{1}{n} \sum_{f=1}^n \mathbf{1}\{L_f \leq \tau \leq U_f\}, \quad \text{Power} = \frac{1}{n} \sum_{f=1}^n \mathbf{1}\{L_f > 0\},$$

so that $\text{MSE} = \text{Bias}^2 + (n-1) \text{SD}^2/n$.

Table S14: Parameter grid of the original eight-design simulation. The crossed parameters form $5 \times 5 \times 2 \times 4 \times 4 \times 2 \times 8 = 12,800$ scenarios, 9,600 of them for the six designs compared in the detailed grid analysis of the detailed analysis; the settings in the lower panel are the same in every scenario.

Parameter	Values	Levels
<i>Crossed</i>		
Treatment effect τ	0.8, 1.0 (primary), 1.5, 2.0, 3.0	5
Incidence configuration	iid; spatially autocorrelated, $\rho_X \in \{0.20, 0.50\}$; Poisson, $\rho_X \in \{0.20, 0.50\}$	5
Neighbor type	Queen (primary), rook (sensitivity)	2
Outcome spatial dependence ρ	0, 0.01, 0.20, 0.50	4
Spillover magnitude γ	0.5, 0.6, 0.7, 0.8	4
Spillover regime	Both arms, control only (Equation S2)	2
Design	The six designs of Table S1 and the two of Table S12	8
<i>Fixed</i>		
Grid	Regular 10×10 grid, $N = 100$ clusters	
Outcome model (Equation S1)	$\beta = 1$, $\sigma = 1$	
Poisson mechanism	Base rate 35 per 100,000; $P = 100,000$ per cluster	
Incidence surfaces per configuration	$K = 10$	
Assignments per surface	$J = 25$ (Checkerboard: one assignment, used in every draw)	
Fits per scenario and estimator	$KJ = 250$, each with its own noise vector	
Estimators	Oracle (primary) and non-oracle (sensitivity), both fit to every simulated outcome	

Table S15: Seed keys. Each key names the kind of draw and the parameters listed here; the incidence mechanism is iid, spatial, or Poisson, and the iid configuration, which has no ρ_X , uses 0 in its place. k indexes the incidence surface and d the design. A draw is shared by every scenario whose key parameters match its own.

Draw	Key parameters	Shared by
Incidence surface X_k	Incidence mechanism, ρ_X , k	Every block, design, and value of τ of the configuration
The $J = 25$ assignments of design d on surface k	Incidence mechanism, ρ_X , neighbor type, ρ , γ , spillover regime, k , d	All five values of τ and both estimators
The 250 noise vectors of a block	Incidence mechanism, ρ_X , neighbor type, ρ , γ , spillover regime	All eight designs, all five values of τ , and both estimators

Surface-level standard errors. For $m \in \{\text{Bias, MSE, Coverage, Power}\}$, let m_k be the same formula applied to the valid fits from surface k (here all 25). With $K = 10$ and $s(a_1, \dots, a_K) = \{(K-1)^{-1} \sum_k (a_k - \bar{a})^2\}^{1/2}$ the sample standard deviation,

$$\text{SE}(m) = \frac{s(m_1, \dots, m_K)}{\sqrt{K}}, \quad m \pm t_{0.975, K-1} \text{SE}(m).$$

Because every fit was valid, each surface contributes 25 fits and $m = K^{-1} \sum_k m_k$. For two designs d and d' in the same block at the same τ , the difference is paired at the surface level: with

$$\Delta_k = m_{d,k} - m_{d',k},$$

$$\text{SE}(m_d - m_{d'}) = \frac{s(\Delta_1, \dots, \Delta_K)}{\sqrt{K}},$$

with the same $t_{0.975, K-1}$ interval.

Standard errors pooled over configurations. For one design, neighbor type, and value of τ , let $\mathcal{B}_c$ be the 32 blocks of configuration c and $m_{b,k}$ the surface- k metric in block b . The configuration-by-surface means and the pooled metric are

$$\bar{m}_{c,k} = \frac{1}{32} \sum_{b \in \mathcal{B}_c} m_{b,k}, \quad \bar{m} = \frac{1}{50} \sum_{c=1}^5 \sum_{k=1}^{10} \bar{m}_{c,k},$$

where $\bar{m}$ equals the mean of the metric over the 160 scenarios, and

$$\text{SE}(\bar{m}) = \frac{s(\{\bar{m}_{c,k}\})}{\sqrt{50}},$$


with s taken over the 50 values $\bar{m}_{c,k}$. These are the standard errors in Table S2. For two designs d and d' adjacent in the mean-MSE order (Formal statistical comparison), let $D_{c,k} = \bar{m}_{d,c,k} - \bar{m}_{d',c,k}$ for $m = \text{MSE}$. The interval for the mean difference is $\bar{D} \pm t_{0.975, 49} s(\{D_{c,k}\})/\sqrt{50}$, and the Wilcoxon signed-rank test, with the normal approximation, is applied to the 50 values $D_{c,k}$.

S6 Supplementary statistical comparisons


This section shows two figures behind results stated in the detailed grid analysis of the detailed analysis.

Figure S10 gives the Nemenyi critical-difference diagrams for the comparison of the six designs at $\tau = 1$, whose results are reported under Formal statistical comparison (queen contiguity) and Rook contiguity in the Results. Each design is placed at its average within-block MSE rank among the six designs over the 160 blocks of its neighbor type. The red arrow shows the critical difference, 0.596 at $\alpha = 0.05$, and a gray bar joins designs whose pairwise Nemenyi comparisons were all not significant: under queen contiguity, Isolation Buffer with 2×2 Blocking, and 2×2 Blocking with Checkerboard and High Incidence Focus; under rook contiguity, Incidence-Guided Saturation Quadrants with Balanced Quartiles, and Isolation Buffer with 2×2 Blocking. Where two designs have nearly equal average ranks (Checkerboard and High Incidence Focus under queen contiguity, Isolation Buffer and 2×2 Blocking under rook contiguity), the dotted leader line to the outer label crosses the neighboring design's label; the positions and labels are correct. Under rook contiguity, τ is not identified for Checkerboard; its estimates are kept and flagged.

Figure S11 gives the coverage of each design across the five values of τ , summarized under Stability across treatment effect sizes. The shaded bands show ± 1 times the average of the scenario-level Monte Carlo standard errors of coverage; they describe the precision of a single scenario's coverage, not of the plotted average. The legend lists the designs in order of pooled mean MSE at $\tau = 1$ for the neighbor type of the panel (Table S2). The vertical axis runs from 0 to 1, so lines close to 0.94 overlap: all six under queen contiguity, between 0.936 and 0.941, and all but Checkerboard under rook contiguity, between 0.937 and 0.942. The lowest line under rook contiguity, which the legend does not mark, is Checkerboard, for which τ is not identified; its estimates are kept and flagged. The five values of τ share their assignments and noise.



(a) Queen contiguity (primary)



(b) Rook contiguity (sensitivity)


Figure S10: Nemenyi critical-difference diagrams for the six designs at $\tau = 1$, pooled over the five incidence configurations. Designs are placed at their average within-block MSE rank, given in parentheses (1 = lowest MSE); gray bars join designs that the Nemenyi comparisons did not separate.

S7 Rook contiguity figures


This section gives the rook-contiguity versions of three figures shown for queen contiguity in the detailed grid analysis of the detailed analysis, as described under Rook contiguity in the Results. Each has the layout of its queen counterpart, but designs are ordered for rook contiguity, so the order differs from that of the queen figure. Under rook contiguity, τ is not identified for Checkerboard ($WZ = 1 - Z$); its estimates are kept and flagged. Figure S12 corresponds to Figure S2.

Figure S13 corresponds to Figure S3. Every design except Checkerboard has coverage close to the nominal level; Checkerboard's coverage, 0.150 pooled over the five incidence configurations (Table S2), is far below it.

Figure S14 corresponds to Figure S4. Nearly all of Checkerboard's MSE is squared bias, because its coefficient on Z estimates $\tau - \gamma$ (Methods).



(a) Queen contiguity (primary)



(b) Rook contiguity (sensitivity)

Figure S11: Coverage of the 95% confidence intervals by design across $\tau \in \{0.8, 1.0, 1.5, 2.0, 3.0\}$, pooled over the five incidence configurations. Each point averages the 160 scenarios of one design at one value of τ . The dashed red line marks the nominal 0.95 level.




Figure S12: Mean squared error of the oracle estimate of τ by design at $\tau = 1$, under rook contiguity; layout as in Figure S2. Checkerboard is flagged (τ not identified).




Figure S13: Coverage of the 95% confidence intervals by design at $\tau = 1$, under rook contiguity; layout as in Figure S3. Checkerboard is flagged (τ not identified).




Figure S14: Squared bias and variance of $\hat{\tau}$ by design at $\tau = 1$ for the Poisson configuration with $\rho_X = 0.20$, under rook contiguity; layout as in Figure S4. Checkerboard is flagged (τ not identified).

S8 Oracle and non-oracle estimates

Table S16 gives the mean squared error, bias, and coverage of the oracle and non-oracle models for each of the six designs, as summarized under Non-oracle sensitivity analysis in the Results.

Table S16: Oracle and non-oracle estimates of τ by design at $\tau = 1$, under queen contiguity (primary) and rook contiguity (sensitivity), pooled over the five incidence configurations. Each entry averages the 160 scenarios of one design and neighbor type. The oracle model has the covariates Z , Spill, and X_k ; the non-oracle model omits Spill. Rows in both panels are in order of oracle MSE under queen contiguity. Under rook contiguity, τ is not identified for Checkerboard ($WZ = 1 - Z$); its estimates are kept and flagged, and the two models coincide.

Design	Oracle			Non-oracle		
	MSE	Bias	Coverage	MSE	Bias	Coverage
<i>Queen contiguity (primary)</i>						
Incidence-Guided	0.091	0.032	0.937	0.087	-0.123	0.810
Saturation Quadrants						
Balanced Quartiles	0.131	0.017	0.936	0.098	-0.176	0.770
Isolation Buffer	0.160	0.003	0.939	0.091	-0.175	0.879
High Incidence Focus	0.245	0.031	0.938	0.219	-0.180	0.896
2×2 Blocking	0.319	0.010	0.939	0.121	-0.244	0.709
Checkerboard	1.087	0.001	0.941	0.108	-0.206	0.741
<i>Rook contiguity (sensitivity)</i>						
Incidence-Guided	0.072	0.018	0.942	0.086	-0.121	0.810
Saturation Quadrants						
Balanced Quartiles	0.084	0.008	0.940	0.097	-0.176	0.772
Isolation Buffer	0.133	0.003	0.939	0.211	-0.400	0.513
High Incidence Focus	0.206	0.018	0.939	0.223	-0.192	0.892
2×2 Blocking	0.128	0.004	0.938	0.128	-0.257	0.691
Checkerboard (flagged)	0.480	-0.652	0.150	0.480	-0.652	0.150

S9 Technical notes

This section explains two results reported in the detailed grid analysis of the detailed analysis: the growth with τ of the bias of Incidence-Guided Saturation Quadrants, and the rise with τ of the coverage of Checkerboard under rook contiguity. The stored results keep neither $\hat{\rho}$ nor the standard error of each fit, so we re-ran every block for the six designs, both neighbor types, and all five values of τ , 2.4 million fits of the oracle model, recording $\hat{\rho}$ and the standard error, and the re-run reproduced the stored results in all 9,600 scenarios. Numbers below that involve $\hat{\rho}$, the standard error, or the terms of Equation S4 come from this re-run; the others come from the stored results.

S9.1 Identification diagnostics

Two combinations of design and neighbor type weaken or remove the information that separates τ from the spillover term. Under rook contiguity, Checkerboard gives $WZ = 1 - Z$ exactly, because every neighbor of a treated cluster is a control and every neighbor of a control is treated. Both spillover covariates then equal $\gamma(1 - Z)$, a linear combination of the intercept and Z , so τ is not

identified: the fitting routine drops the aliased spillover column, and the coefficient on Z estimates $\tau - \gamma$. the earlier small-grid study noted the same consequence of exact collinearity between treatment and spillover exposure, $\hat{\beta} \approx \beta - \psi$ in its notation. The non-oracle model has no spillover column to drop, but its Checkerboard estimates under rook contiguity have the same defect. Under queen contiguity, each of the 64 interior clusters has four treated and four control neighbors, so $WZ = 1/2$ at every interior cluster. Under the control-only regime the spillover covariate then equals $(\gamma/2)(1 - Z)$ at those clusters, a linear combination of the intercept and Z , and only the 36 clusters on the grid boundary carry information that separates it from the intercept and Z ; τ is identified, but imprecisely. Under the both-arms regime it is the constant $\gamma/2$ at interior clusters, confounded with the intercept but not with Z . Isolation Buffer under rook contiguity can also, rarely, draw exactly the checkerboard pattern, a maximal independent set that equals one color class of the grid. This happened in 4 of its 40,000 rook draws and affected one fit in each of 20 scenarios, 4 at each value of τ . We flag every fit in which the fitting routine reports an aliased column and every scenario in which some draw has $\text{rank}([1, Z, WZ]) < 3$. Aliased estimates are kept in all summaries. In the exhibits, we flag Checkerboard under rook contiguity, where every fit is aliased; the Isolation Buffer scenarios, each with one aliased fit among its 250, are noted here and carry no flag in the exhibits.

S9.2 Estimation error and its dependence on τ

For one fit (k, j) , write Y , Z , and ϵ for Y_{kj} , Z_{kj} , and ϵ_{kj} (Simulation design, resampling, and metrics). Let $V = [1, Z, S(Z), X_k]$ be the 100×4 covariate matrix of the oracle model, so that Equation S1 reads $(I - \rho W)Y = V\theta + \epsilon$ with $\theta = (0, \tau, 1, \beta)^\top$. Let $A = (I - \rho W)^{-1}$ and $V^+ = (V^\top V)^{-1}V^\top$, and for a vector v write $[V^+v]_Z$ for the coefficient on Z when v is regressed on the columns of V by least squares. Given $\hat{\rho}$, the estimated coefficients are $V^+(I - \hat{\rho}W)Y$ (Estimation). Because $(I - \hat{\rho}W)Y = V\theta + \epsilon + (\rho - \hat{\rho})WY$ and $WY = WA(V\theta + \epsilon)$,

$$\hat{\tau} - \tau = [V^+\epsilon]_Z + (\rho - \hat{\rho})h + \tau(\rho - \hat{\rho})g, \quad (\text{S4})$$

where

$$g = [V^+WAZ]_Z, \quad h = [V^+WA\{S(Z) + \beta X_k + \epsilon\}]_Z.$$

The first term has mean zero given Z and X_k . The last term carries a factor of τ . It is present exactly when $g \neq 0$, where g is the coefficient on Z when the lagged treatment $WAZ = WZ + \rho W^2Z + \rho^2 W^3Z + \dots$ is regressed on the model's own covariates, and its sign and size are set by the error in $\hat{\rho}$. The concentrated likelihood depends on Y only through the residual sum of squares of $(I - \hat{\rho}W)Y$ on V , and the part of that residual that involves τ is $\tau(\rho - \hat{\rho})$ times the residual of WAZ on V . If WAZ lies in the column space of V , that part vanishes for every $\hat{\rho}$, so $\hat{\rho}$ is the same at every value of τ ; if in addition $g = 0$, then $\hat{\tau} - \tau$ is the same at every value of τ . Equation S4 held on every re-run fit, with a largest residual of 4.2×10^{-13} ; for Checkerboard under rook contiguity it held in the reduced form given below.

Two consequences of Equation S4 explain the results for τ reported in the detailed grid analysis.

First, for designs whose treated clusters have treated neighbors, the last term gives a small positive bias that grows with τ . Under queen contiguity, pooled over the 160 blocks, the bias of Incidence-Guided Saturation Quadrants rose from 0.029 at $\tau = 0.8$ to 0.047 at $\tau = 3$ (Stability across treatment effect sizes), while that of Isolation Buffer stayed between 0.002 and 0.004. Almost all of the rise is the last term of Equation S4. In our fits, $\hat{\rho}$ was biased downward by about 0.06 on average (for Incidence-Guided Saturation Quadrants, by 0.071 at $\tau = 0.8$ and 0.054 at $\tau = 3$). We observed this downward bias in these simulations; we did not derive it or check it against a

published finite-sample result. With $\rho - \hat{\rho}$ positive on average, the last term is positive and grows with τ for any design with $g > 0$. For Incidence-Guided Saturation Quadrants, g averaged 0.371 over the queen fits, and the last term contributed 0.021 of the bias at $\tau = 0.8$ and 0.043 at $\tau = 3$; the first term contributed 0.001 at both values, and the second 0.007 and 0.002. For Isolation Buffer, g averaged 0.020, and the last term contributed 0.001 and 0.004. Most of the difference in g comes from the control-only regime. Under the both-arms regime with $\rho = 0$, the lagged treatment is $WZ = S(Z)/\gamma$, one of the covariates, so $g = 0$ for every design. Under the control-only regime, $WZ = S(Z)/\gamma + Z \circ WZ$, and at $\rho = 0$, g is close to the average share of treated neighbors among treated clusters. That share is 0.57 for Incidence-Guided Saturation Quadrants, whose treated clusters are concentrated in the high-saturation quadrants, and exactly 0 for Isolation Buffer, whose treated clusters have no treated neighbors. For Isolation Buffer, $Z \circ WZ = 0$, so the two regimes define the same spillover covariate and the same g , which is 0 at $\rho = 0$ and about 0.06 at $\rho = 0.5$; at $\rho = 0$, its $\hat{\rho}$ and its estimation error were the same at every value of τ , up to the optimizer’s tolerance.

The same term explains why the control-only MSE of Incidence-Guided Saturation Quadrants rose from 0.131 at $\tau = 0.8$ to 0.182 at $\tau = 3$ while that of Isolation Buffer did not. Under that regime, averaged over the 80 queen blocks, the variance of the last term of Equation S4 rose from 0.008 to 0.061 for Incidence-Guided Saturation Quadrants, while the variance of the rest of the error changed little (0.119 and 0.115). The variance of the last term grew more slowly than τ^2 because $\hat{\rho}$ became less variable as τ grew. For Isolation Buffer, the variance of the last term stayed below 0.0002. The rise in bias slowed at larger values of τ for a related reason: the part of the likelihood that involves τ sharpens it around the true ρ as τ grows, so the downward bias of $\hat{\rho}$ shrank.

Second, the decomposition explains why the coverage of Checkerboard under rook contiguity rose with τ . Under rook contiguity, Checkerboard’s coverage rose from 0.138 at $\tau = 0.8$ to 0.442 at $\tau = 3$ (Stability across treatment effect sizes), although its coefficient on Z estimates $\tau - \gamma$ at every value of τ . For this design $WZ = 1 - Z$, and because W is row-standardized, $W^2Z = W(1 - Z) = Z$. The powers of W applied to Z therefore alternate between $1 - Z$ and Z , and $WAZ = \{(1 - Z) + \rho Z\}/(1 - \rho^2)$ lies in the span of the intercept and Z . The derivation of Equation S4 applies with the aliased spillover column dropped, $V = [1, Z, X_k]$ and $\theta = (\gamma, \tau - \gamma, \beta)^\top$, so that $\tau - \gamma$ takes the place of τ , and $g = -1/(1 + \rho)$. Two consequences follow: $\hat{\rho}$ is the same at every value of τ , and

$$\hat{\tau} - \tau = -\gamma + \delta + (\tau - \gamma) \frac{\hat{\rho} - \rho}{1 + \rho},$$

where δ collects the first two terms of that decomposition and, like $\hat{\rho}$, does not depend on τ . The error in $\hat{\rho}$ is therefore multiplied by $\tau - \gamma$, which grows with τ . The mean standard deviation of $\hat{\tau}$ rose from 0.202 at $\tau = 0.8$ to 0.327 at $\tau = 3$, and the standard error tracked it, rising on average from 0.199 to 0.322. With the bias held near $-\gamma$, wider intervals covered τ more often. The rise was largest at $\rho = 0$ (coverage 0.145 to 0.541) and smallest at $\rho = 0.5$ (0.130 to 0.265), consistent with the factor $1/(1 + \rho)$. The same term raised the MSE from 0.477 to 0.567. About three quarters of that increase was variance, $(n - 1)SD^2/n$ averaged over scenarios (Section S5), which rose from 0.041 to 0.109. The other quarter was squared bias, which rose from 0.436 to 0.458 because $\hat{\rho}$ was biased slightly downward, moving the mean bias from -0.650 to -0.667 . The higher coverage at large τ reflects wider intervals around the wrong quantity, not a recovery of τ .

S10 Observed-incidence application detail

The primary application numbers are traced to the production performance and allocation-metric CSVs in the versioned application output directory and the separate `tail_confirmation/` outputs. Yearly exhibits are in `exhibits/`; annual tables average the six queen parameter pairs equally, without pooled Monte Carlo intervals. Content-keyed Source.IDs identify shared computational sources. Manifests cover configuration, code, input and cached-weight hashes, packages, BLAS and RNG settings; changed computation refuses checkpoint reuse. Independent verification recomputes performance from every source cache, checks annual budgets and named ordering, and checks all refined fit/allocation prefixes. Restricted county inputs and individual-fit caches remain excluded from the public repository.

For stochastic supports the prespecified refinement tiers allow $J = 100, 200, 400$ and $R = 100, 200, 400, 1000, 2000, 4000$; the component contributing most to normalized precision short-fall determines whether allocations or outcomes are extended. Main stochastic sources passed at $J = 100, R = 100$. For singleton supports only outcomes are extended. If an allocation occurs n_a times among J draws, its risk estimate enters the mean with weight n_a/J . Conditional-noise variance therefore involves squared weights $(n_a/J)^2$, not n_a/J^2 as if duplicate estimates were independent. Bias and coverage use the same frequency/covariance accounting. Joint precision includes assignment variation and conditional outcome noise; zero sampled variation is distinguished from proven singleton support. Approximate setting-specific design-difference MCSEs use independent design streams, while cross-setting/model averages do not assume independence of shared allocations.

Estimated risk variance is reported raw and with its outcome Monte Carlo noise contribution subtracted. Signed corrected estimates are retained rather than silently truncated; none is negative in the NC refinement. Quantiles and worst-decile means use noisy estimated conditional risks. For selection diagnostics, one half of outcomes ranks assignments, and the opposite half estimates risk for the selected worst decile; directions are then reversed. Their Jaccard overlap describes sampled tail membership, not a confidence interval or population tail probability. A singleton’s overlap of one is vacuous.

The frozen regional criterion is $C = 3s_P + 1.5s_K + 1.5s_D + c/d_{\max}$, where s_P , s_K and s_D are the standard deviations of regional population, cluster and county shares. Here c averages within-region distances to regional centers and $d_{\max}$ is the maximum pairwise distance between the projected cluster coordinates. Candidates use standardized projected coordinates and are selected without outcome performance. Frozen block sizes range from one to eight, with one singleton per year; yearly total treatments are 28,31,27,29. Table S17 documents alternative summaries before performance evaluation.

Table S17: Regional ordering from highest to lowest summary, using the frozen region IDs in Figure S15. Primary mean ranks weight clusters equally; pooled regional rates weight cluster rates by population.

Year	Mean cluster rank	Mean cluster rate	Pooled regional rate
2018	2,1,3,4	2,4,1,3	4,3,2,1
2019	4,2,3,1	4,2,3,1	4,3,1,2
2020	2,4,1,3	2,4,1,3	4,3,2,1
2021	2,1,4,3	2,4,1,3	4,3,1,2

The following annual tables retain bias, coverage and both budget measures for all designs.

One frozen region partition for all four years
 Selected by geography and balance before outcome performance is evaluated




Figure S15: Frozen four-region partition used by both saturation strategies in all years, chosen from 250 geographic candidates before outcome evaluation.

Reproducible example allocations on the NC geography

2018, queen graph, draw 1. These maps illustrate rules; they do not compare performance.




Figure S16: One illustrative queen assignment from each application strategy using 2018 observed incidence and frozen geography. These illustrative masks explain adaptation; they are not selected by observed performance.

More granular setting-specific results, precision components, residual treatment information and individual warnings/failures tables are supplied as the reproducible aggregate CSV outputs. Zero warning/failure rows are explicit verified results, not evidence that diagnostics were skipped.

Table S22 reports guided saturation's matched annual mean-MSE ratios. Rook means use only the four shared parameter corners. Baseline/X-adjusted and regional-summary means use the six queen pairs. These are matched contrasts, not ratios of unmatched pooled analyses.




Figure S17: Annual population shares treated under primary queen assignments. Incidence-independent supports share performance but annual population shares are recalculated. Population aged 18–64 is geographic reach rather than education-participant enrollment.




Figure S18: Primary queen setting-specific mean MSE ratios to year-matched SRS. Parameters and spillover regimes are kept separate; this plot does not pool neighbor definitions.

Table S18: Primary queen 2018 application performance. Equal weights over six parameter pairs within each regime; bias, MSE and coverage are Monte Carlo estimates. Treatment count and population share describe the allocation budget.

Design	MSE	Bias	Coverage	Treated	Pop. %
<i>Both arms</i>					
Graph Checkerboard	0.21591	-0.0074	0.936	27.00	50.4
High Incidence Focus	0.07172	0.0160	0.934	29.00	22.2
Plain saturation	0.07518	0.0208	0.935	29.03	49.2
Isolation Buffer	0.28591	0.0088	0.932	15.89	26.7
Spatial blocking	0.07609	0.0099	0.933	28.00	50.0
Balanced Quartiles	0.07610	0.0170	0.933	29.00	48.5
Balanced Halves	0.07458	0.0158	0.934	29.00	50.0
Guided saturation	0.07428	0.0174	0.935	29.00	53.2
SRS	0.07638	0.0135	0.932	29.00	49.5
<i>Control only</i>					
Graph Checkerboard	0.92036	0.0100	0.933	27.00	50.4
High Incidence Focus	0.21426	0.0237	0.932	29.00	22.2
Plain saturation	0.22329	0.0464	0.926	29.03	49.2
Isolation Buffer	0.28575	0.0106	0.931	15.89	26.7
Spatial blocking	0.30622	0.0281	0.929	28.00	50.0
Balanced Quartiles	0.29352	0.0315	0.929	29.00	48.5
Balanced Halves	0.29101	0.0395	0.930	29.00	50.0
Guided saturation	0.24285	0.0464	0.928	29.00	53.2
SRS	0.30761	0.0378	0.929	29.00	49.5

Table S19: Primary queen 2019 application performance. Equal weights over six parameter pairs within each regime; bias, MSE and coverage are Monte Carlo estimates. Treatment count and population share describe the allocation budget.

Design	MSE	Bias	Coverage	Treated	Pop. %
<i>Both arms</i>					
Graph Checkerboard	0.21591	-0.0074	0.936	27.00	50.4
High Incidence Focus	0.07180	0.0063	0.936	29.00	23.5
Plain saturation	0.07518	0.0208	0.935	29.03	49.2
Isolation Buffer	0.28591	0.0088	0.932	15.89	26.7
Spatial blocking	0.07808	0.0143	0.930	31.00	55.1
Balanced Quartiles	0.07624	0.0175	0.933	29.00	49.7
Balanced Halves	0.07517	0.0167	0.934	29.00	50.5
Guided saturation	0.07628	0.0233	0.935	29.00	45.1
SRS	0.07638	0.0135	0.932	29.00	49.5
<i>Control only</i>					
Graph Checkerboard	0.92036	0.0100	0.933	27.00	50.4
High Incidence Focus	0.43019	0.0227	0.931	29.00	23.5
Plain saturation	0.22329	0.0464	0.926	29.03	49.2
Isolation Buffer	0.28575	0.0106	0.931	15.89	26.7
Spatial blocking	0.43630	0.0314	0.928	31.00	55.1
Balanced Quartiles	0.30616	0.0367	0.929	29.00	49.7
Balanced Halves	0.29955	0.0360	0.930	29.00	50.5
Guided saturation	0.21565	0.0522	0.926	29.00	45.1
SRS	0.30761	0.0378	0.929	29.00	49.5

Table S20: Primary queen 2020 application performance. Equal weights over six parameter pairs within each regime; bias, MSE and coverage are Monte Carlo estimates. Treatment count and population share describe the allocation budget.

Design	MSE	Bias	Coverage	Treated	Pop. %
<i>Both arms</i>					
Graph Checkerboard	0.21591	-0.0074	0.936	27.00	50.4
High Incidence Focus	0.07292	0.0094	0.931	29.00	23.3
Plain saturation	0.07518	0.0208	0.935	29.03	49.1
Isolation Buffer	0.28591	0.0088	0.932	15.89	26.7
Spatial blocking	0.07764	0.0122	0.933	27.00	48.0
Balanced Quartiles	0.07530	0.0158	0.933	29.00	50.6
Balanced Halves	0.07466	0.0150	0.934	29.00	49.2
Guided saturation	0.07682	0.0209	0.933	29.00	47.9
SRS	0.07638	0.0135	0.932	29.00	49.5
<i>Control only</i>					
Graph Checkerboard	0.92036	0.0100	0.933	27.00	50.4
High Incidence Focus	0.25918	0.0327	0.930	29.00	23.3
Plain saturation	0.22329	0.0464	0.926	29.03	49.1
Isolation Buffer	0.28575	0.0106	0.931	15.89	26.7
Spatial blocking	0.31848	0.0281	0.930	27.00	48.0
Balanced Quartiles	0.28932	0.0356	0.927	29.00	50.6
Balanced Halves	0.30610	0.0372	0.930	29.00	49.2
Guided saturation	0.22029	0.0474	0.926	29.00	47.9
SRS	0.30761	0.0378	0.929	29.00	49.5

Table S21: Primary queen 2021 application performance. Equal weights over six parameter pairs within each regime; bias, MSE and coverage are Monte Carlo estimates. Treatment count and population share describe the allocation budget.

Design	MSE	Bias	Coverage	Treated	Pop. %
<i>Both arms</i>					
Graph Checkerboard	0.21591	-0.0074	0.936	27.00	50.7
High Incidence Focus	0.07031	0.0104	0.934	29.00	22.4
Plain saturation	0.07518	0.0208	0.935	29.03	49.1
Isolation Buffer	0.28591	0.0088	0.932	15.89	26.7
Spatial blocking	0.07736	0.0107	0.931	29.00	50.3
Balanced Quartiles	0.07453	0.0175	0.934	29.00	49.8
Balanced Halves	0.07584	0.0160	0.932	29.00	49.4
Guided saturation	0.07450	0.0197	0.936	29.00	51.8
SRS	0.07638	0.0135	0.932	29.00	49.5
<i>Control only</i>					
Graph Checkerboard	0.92036	0.0100	0.933	27.00	50.7
High Incidence Focus	0.35318	0.0296	0.929	29.00	22.4
Plain saturation	0.22329	0.0464	0.926	29.03	49.1
Isolation Buffer	0.28575	0.0106	0.931	15.89	26.7
Spatial blocking	0.36168	0.0289	0.928	29.00	50.3
Balanced Quartiles	0.28439	0.0354	0.929	29.00	49.8
Balanced Halves	0.29089	0.0340	0.929	29.00	49.4
Guided saturation	0.23732	0.0451	0.926	29.00	51.8
SRS	0.30761	0.0378	0.929	29.00	49.5

Table S22: Guided saturation sensitivity MSE / matched primary MSE by year and regime. Baseline adds rank-scaled X to both generation and fit; rook uses matched corners. Regional-summary alternatives preserve frozen regions.

Sensitivity	Regime	2018	2019	2020	2021
Baseline / adjusted	Both	1.016	1.030	1.015	1.035
Baseline / adjusted	Control	1.026	1.023	1.042	1.040
Rook corners	Both	1.005	0.998	0.985	0.998
Rook corners	Control	0.967	0.969	0.971	1.036
Mean regional rate	Both	1.034	1.000	1.000	1.031
Mean regional rate	Control	0.907	1.000	1.000	0.928
Pooled regional rate	Both	1.015	0.979	0.981	1.003
Pooled regional rate	Control	0.962	1.133	1.061	1.030

References

- Andrew Walther, Tonya Van Deinse, and Feng-Chang Lin. Spatial cluster randomized trials - sampling design with spillover effects & spatial dependence. *BMC Medical Research Methodology*, 2026. doi: 10.1186/s12874-026-02968-0. URL <https://link.springer.com/article/10.1186/s12874-026-02968-0>. Published online 26 August 2026.
- Mojtaba Mirzaei, Golsa Joodi, Brittany Bogle, Sarah Chen, and Ross J. Simpson. Years of Life and Productivity Loss Because of Adult Sudden Unexpected Death in the United States. *Medical Care*, 57(7):498–502, 2019. ISSN 0025-7079. doi: 10.1097/MLR.0000000000001129. URL <https://journals.lww.com/00005650-201907000-00003>.
- Roger Bivand and Gianfranco Piras. Comparing Implementations of Estimation Methods for Spatial Econometrics. *Journal of Statistical Software*, 63(18), 2015. ISSN 1548-7660. doi: 10.18637/jss.v063.i18. URL <http://www.jstatsoft.org/v63/i18/>.
- Keith Ord. Estimation Methods for Models of Spatial Interaction. *Journal of the American Statistical Association*, 70(349):120–126, 1975. ISSN 0162-1459, 1537-274X. doi: 10.1080/01621459.1975.10480272. URL <http://www.tandfonline.com/doi/abs/10.1080/01621459.1975.10480272>.
- Milton Friedman. The use of ranks to avoid the assumption of normality implicit in the analysis of variance. *Journal of the American Statistical Association*, 32(200):675–701, 1937. doi: 10.1080/01621459.1937.10503522.
- Peter Bjorn Nemenyi. *Distribution-Free Multiple Comparisons*. PhD thesis, Princeton University, Princeton, NJ, 1963.
- Janez Demšar. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research*, 7:1–30, 2006. URL <https://jmlr.org/papers/v7/demsar06a.html>.
- Frank Wilcoxon. Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83, 1945. doi: 10.2307/3001968.
- Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979. URL <https://www.jstor.org/stable/4615733>.
- Ashkan Habib. Temporal trends and social determinants of sudden death mortality in North Carolina. Master’s paper, Department of Biostatistics, University of North Carolina at Chapel Hill, Chapel Hill, NC, April 2026. Unpublished.